

Distributed Systems Architectures & Processes

Olaf Landsiedel

Notes

- Labs
 - Please submit in teams of two!

Last Time

- Communication
 - Tour de Protocol Stack
 - Hour glass model

Today

- Architectures
 - Design choices for distributed systems
- Processes
 - Refresh your knowledge on concurrency etc.

Part I

Architecture

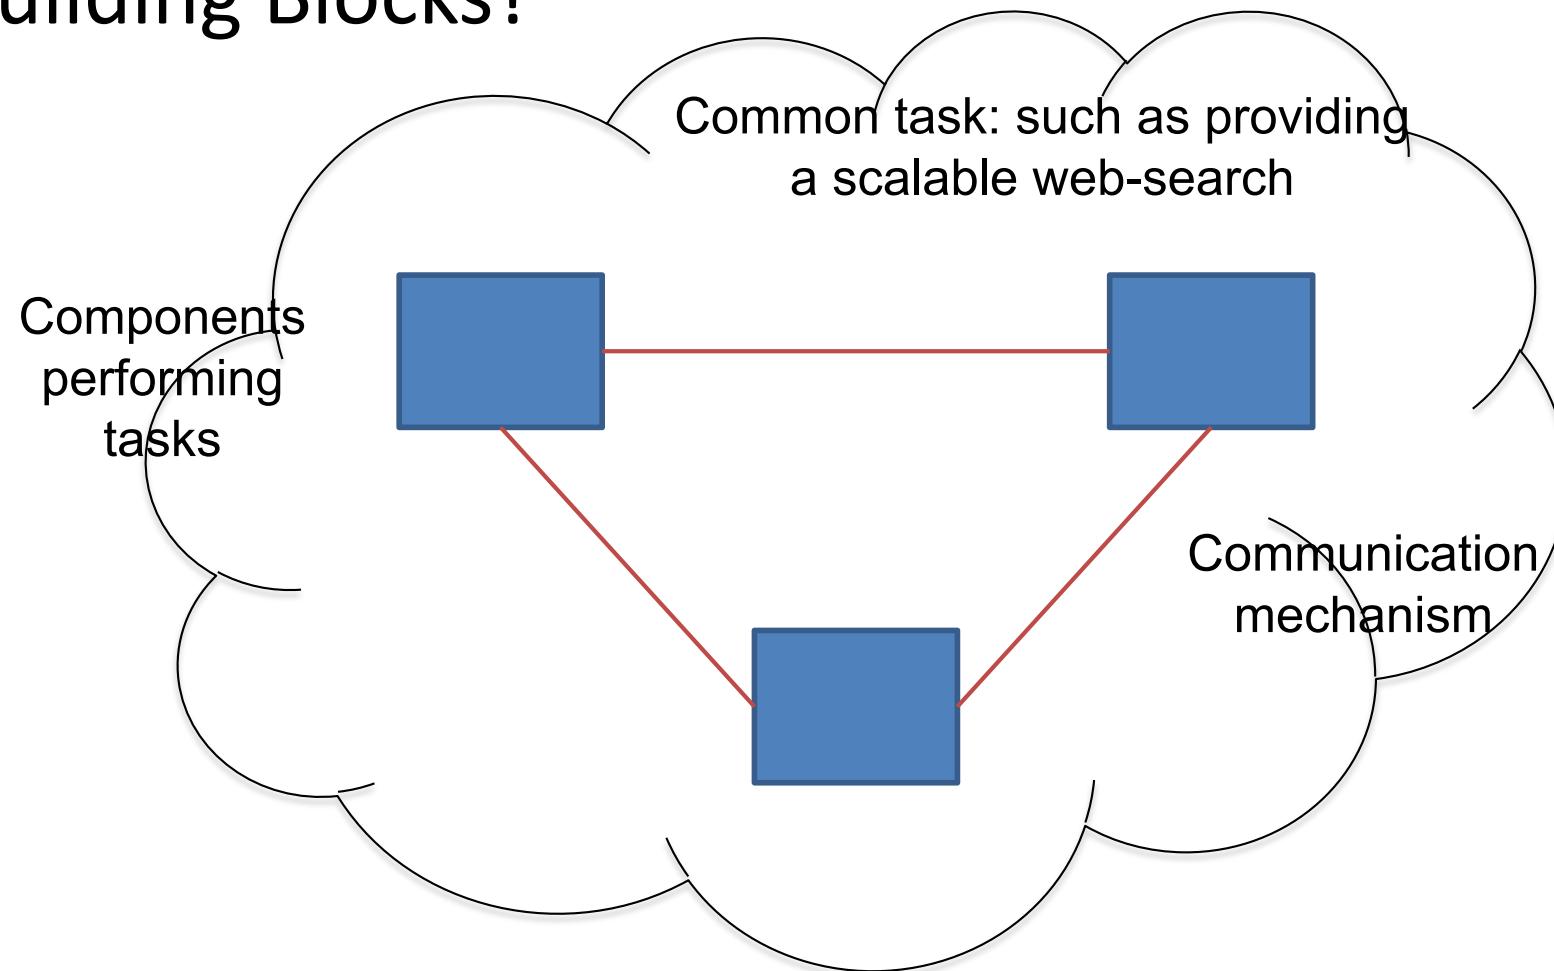
A Distributed System

- A Distributed System is characterized by?
 - Multiple devices
 - Connected by a network
 - Cooperating on some task
- You know when you have one ...
... when the failure of a computer you've never heard of stops you from getting any work done

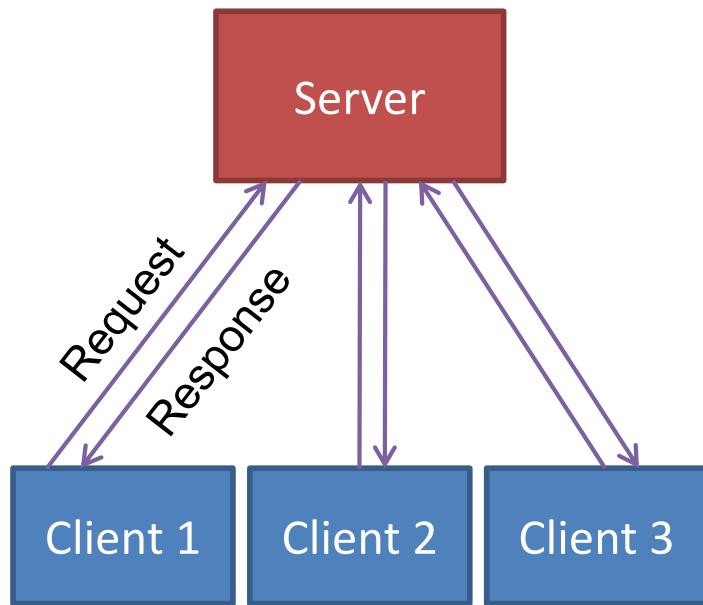
(L. Lamport)

A Distributed System

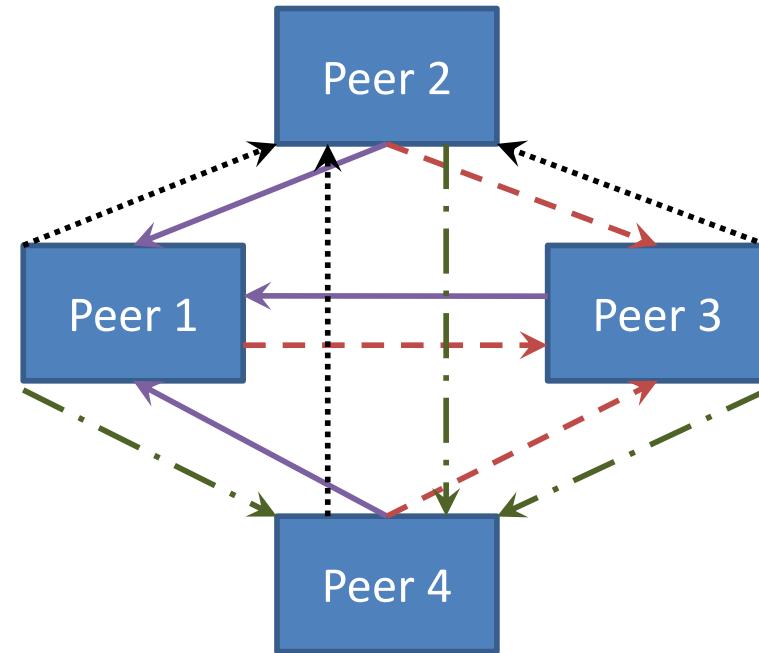
- Building Blocks?



Building Distributed Systems: Design Choices



Client – Server Model



Peer-To-Peer Model

Pros and Cons?

Classification of Distributed Systems

- What are the entities that are communicating in a DS?
 - a) **Communicating entities**
- How do the entities communicate?
 - b) **Communication paradigms**
- What roles and responsibilities do they have?
 - c) **Roles and responsibilities**
- How are they mapped to the physical distributed infrastructure?
 - d) **Placement of entities**

Classification of Distributed Systems

- What are the entities that are communicating in a DS?
 - a) **Communicating entities**
- How do the entities communicate?
 - b) Communication paradigms
- What roles and responsibilities do they have?
 - c) Roles and responsibilities
- How are they mapped to the physical distributed infrastructure?
 - d) Placement of entities

Communicating Entities

- What entities are communicating in a DS?
 - System-oriented entities
 - Processes
 - Threads
 - Nodes
 - Problem-oriented entities
 - Objects (in *object-oriented programming* based approaches)

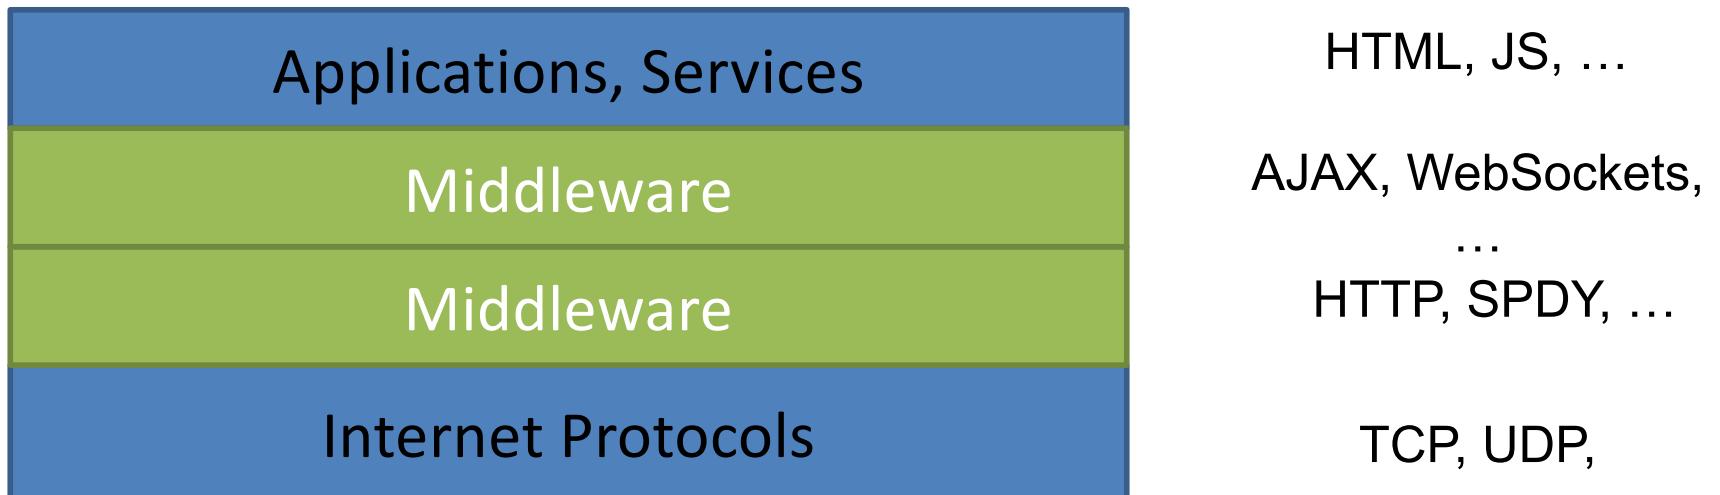
Classification of Distributed Systems

- What are the entities that are communicating in a DS?
 - a) Communicating entities
- How do the entities communicate?
 - b) **Communication paradigms**
- What roles and responsibilities do they have?
 - c) Roles and responsibilities
- How are they mapped to the physical distributed infrastructure?
 - d) Placement of entities

Communication Paradigm

- Middleware
 - Service “Library”
 - On server and client

Example



Problems

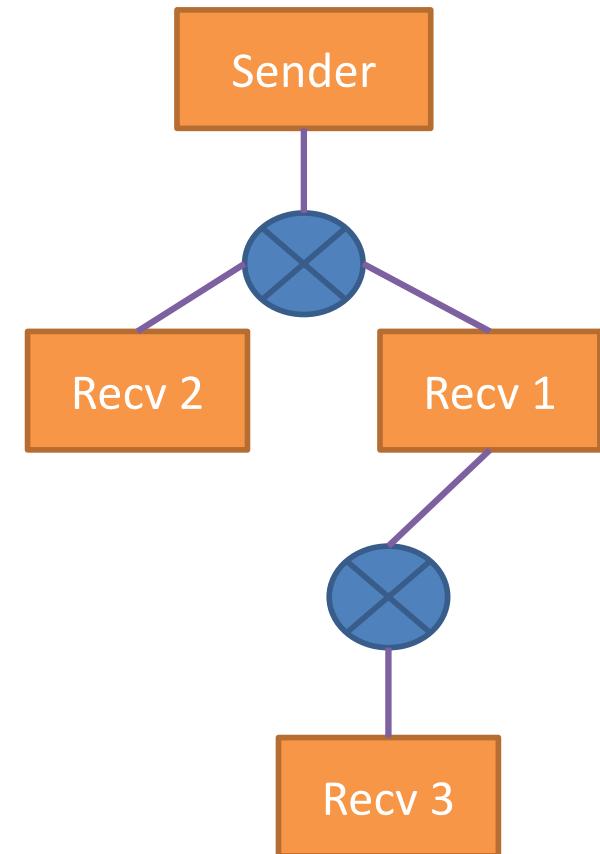
- **Space Coupling:**
 - Where the procedure resides should be known in advance
- **Time Coupling:**
 - On the receiver, a process should be explicitly waiting to accept requests for procedure calls
- Next: Indirect communication
 - Aiming to fix this

Indirect Communication Paradigm

- Indirect communication uses middleware to:
 - Provide one-to-many communication
 - Some mechanisms eliminate space and time coupling
 - Sender and receiver do not need to know each other's identities
 - Sender and receiver need not be explicitly listening to communicate
- Approach used: **Indirection**
 - Sender → **A middle-man** → Receiver
- Types of indirect communication
 1. Group communication
 2. Publish-subscribe
 3. Message queues

1. Group Communication

- One-to-many communication
 - Multicast communication
- Abstraction of a group
 - Group is represented in the system by a *groupId*
 - Recipients join the group
 - A sender sends a message to the group which is received by all the recipients

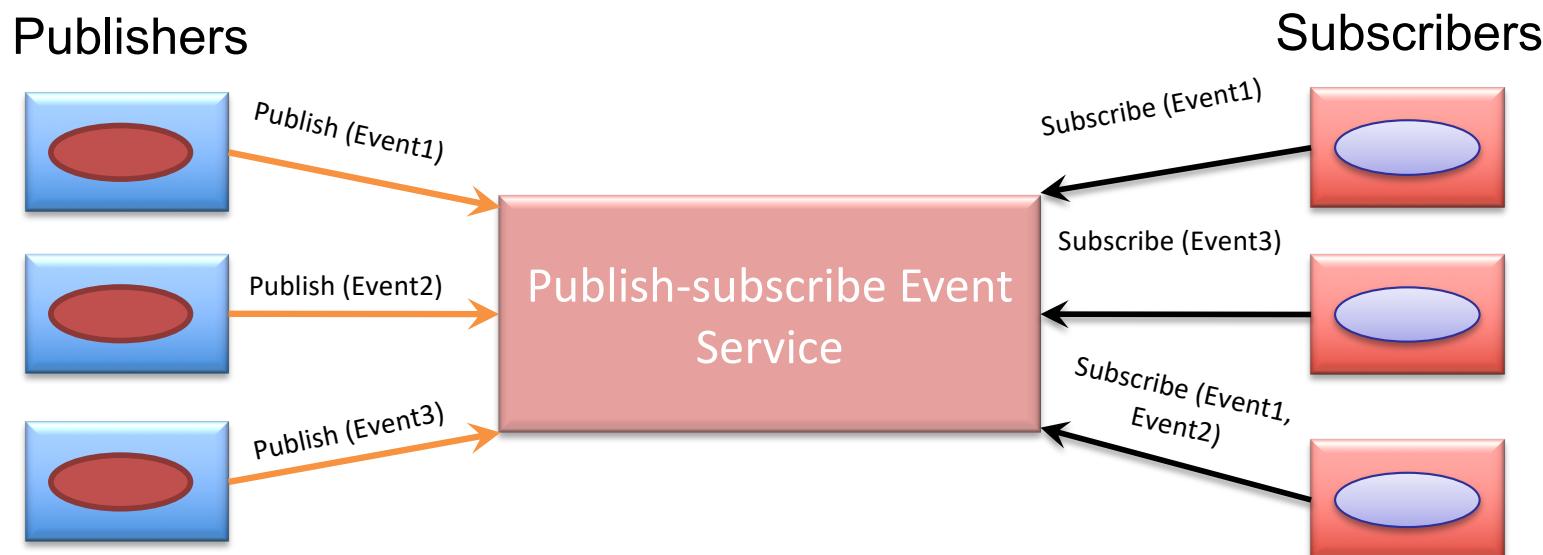


1. Group Communication (cont'd)

- Services provided by middleware
 - Group membership
 - Handling the failure of one or more group members
- Advantages?
 - Enables one-to-many communication
 - Efficient use of bandwidth
 - Identity of the group members need not be available at all nodes
- Disadvantages
 - Time coupling

2. Publish-Subscribe

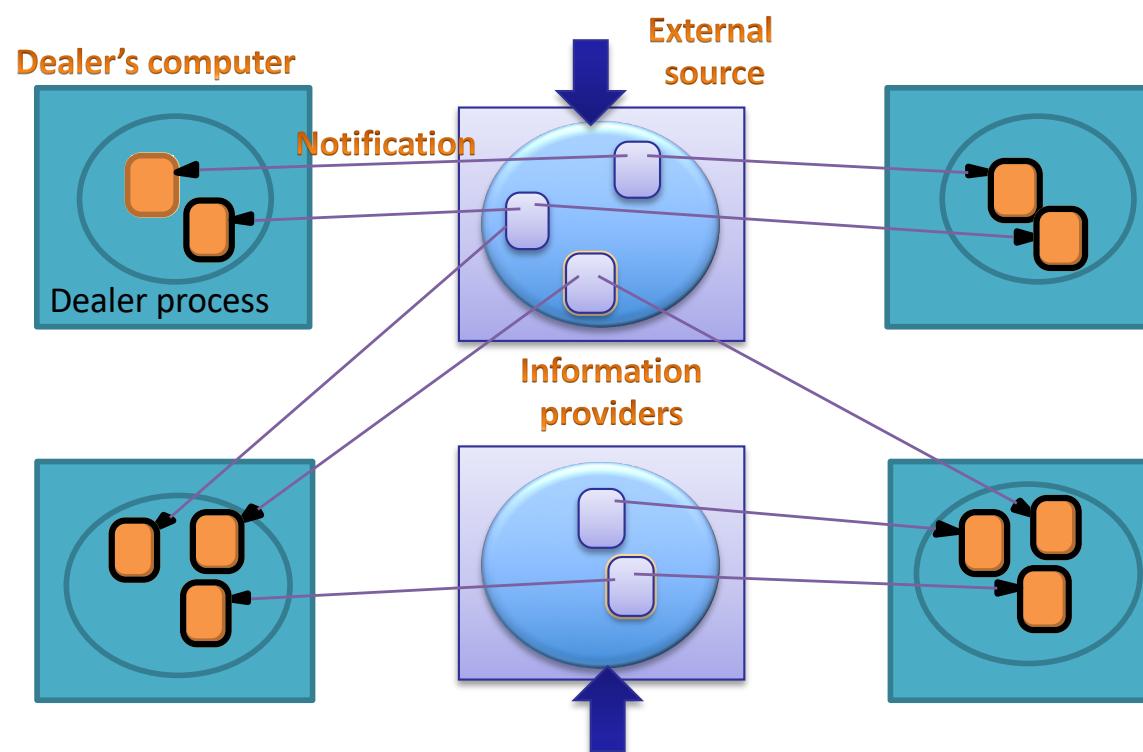
- An event-based communication mechanism
 - Publishers publish events to an event service
 - Subscribers express interest in particular events



- Large number of producers distribute information to large number of consumers

2. Publish-Subscribe (cont'd)

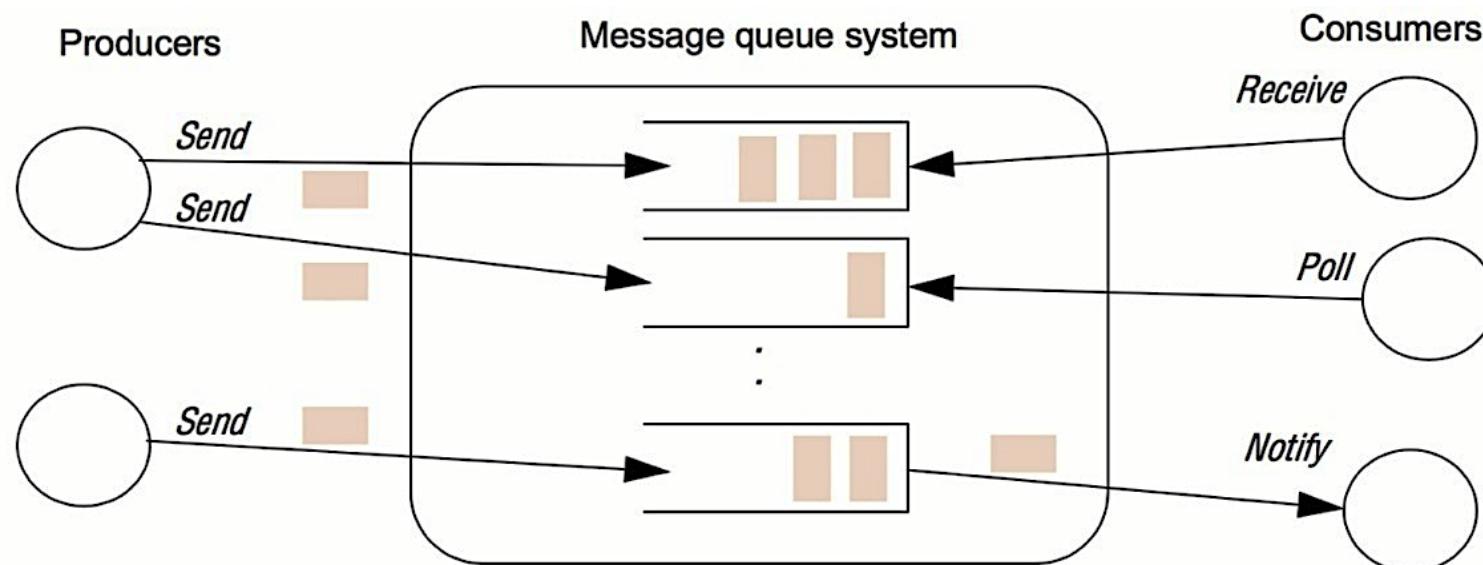
Example: Financial trading



When is publish subscribe good?

3. Message Queues

- A refinement of Publish-Subscribe where
 - Producers deposit the messages in a queue
 - Messages are delivered to consumers through different methods
 - Queue takes care of ensuring message delivery
- Advantages
 - Enables space decoupling
 - Enables time decoupling



Last Time

- Tour de Protocol Stack, Part II
- Architectures of Distributed Systems

Recap: Communication Entities and Paradigms

<i>Communicating entities (what is communicating)</i>		<i>Communication Paradigms (how they communicate)</i>		
<i>System-oriented</i>	<i>Problem-oriented</i>	<i>“Raw”</i>	<i>Middleware</i>	<i>Indirect Communication</i>
<ul style="list-style-type: none">• Nodes• Processes• Threads	<ul style="list-style-type: none">• Objects	<ul style="list-style-type: none">• Sockets	<ul style="list-style-type: none">• “Library”	<ul style="list-style-type: none">• Group communication• Publish-subscribe• Message queues

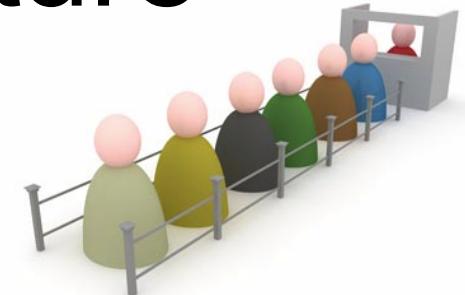
Classification of Distributed Systems

- What are the entities that are communicating in a DS?
 - a) Communicating entities
- How do the entities communicate?
 - b) Communication paradigms
- What roles and responsibilities do they have?
 - c) **Roles and responsibilities**
- How are they mapped to the physical distributed infrastructure?
 - d) Placement of entities

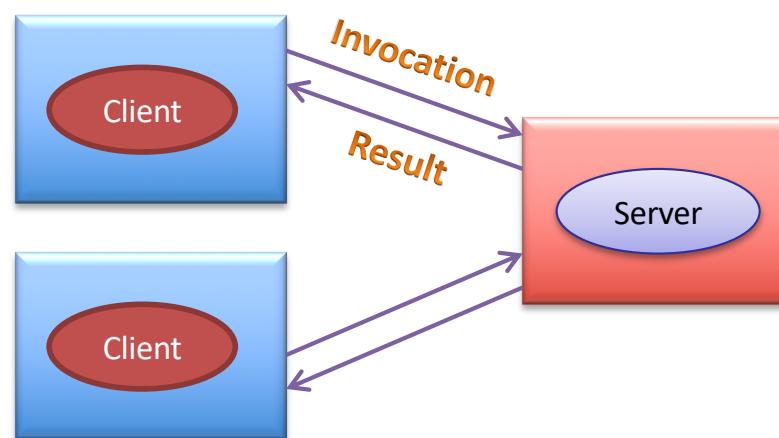
Roles and Responsibilities

- In DS, communicating entities take on roles to perform tasks
- Roles are fundamental in establishing overall architecture
 - Question: Does your smart-phone perform the same role as Google Search Server?
- We classify DS architectures into two types based on the roles and responsibilities of the entities
 - Client-Server
 - Peer-to-Peer

Client-Server Architecture



- Approach:
 - Server provides a service that is needed by a client
 - Client requests to a server (invocation), the server serves (result)
- Widely used in many systems
 - e.g., DNS, Web-servers

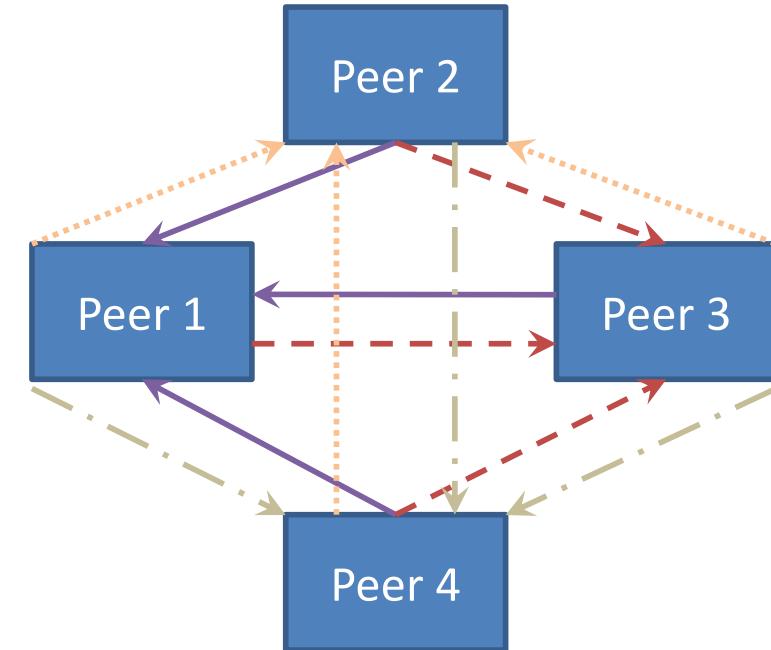


Client-Server Architecture: Pros and Cons

- Advantages?
 - Simplicity and centralized control
 - Computation-heavy processing can be offloaded to a powerful server, Clients can be “thin”
- Disadvantages?
 - Single-point of failure at server
 - Scalability

Peer to Peer (P2P) Architecture

- In P2P, roles of all entities are identical
 - All nodes are peers
 - Peers are equally privileged participants in the application
- Examples
 - Napster, BitTorrent, Skype, Spotify, ...
(see also next slides)



Peer-To-Peer Model

Peer to Peer (P2P) Architecture

- Advantages
 - Scalable design
 - Limited cost for service provider
- Disadvantages
 - Limited control
 - Free riding
 - Churn: nodes coming and going
 - Bandwidth
 - ...
- For this reason:
 - Skype, Spotify, etc. moved away from P2P model
 - Used only during the beginning of the company
 - When they had limited financial resources

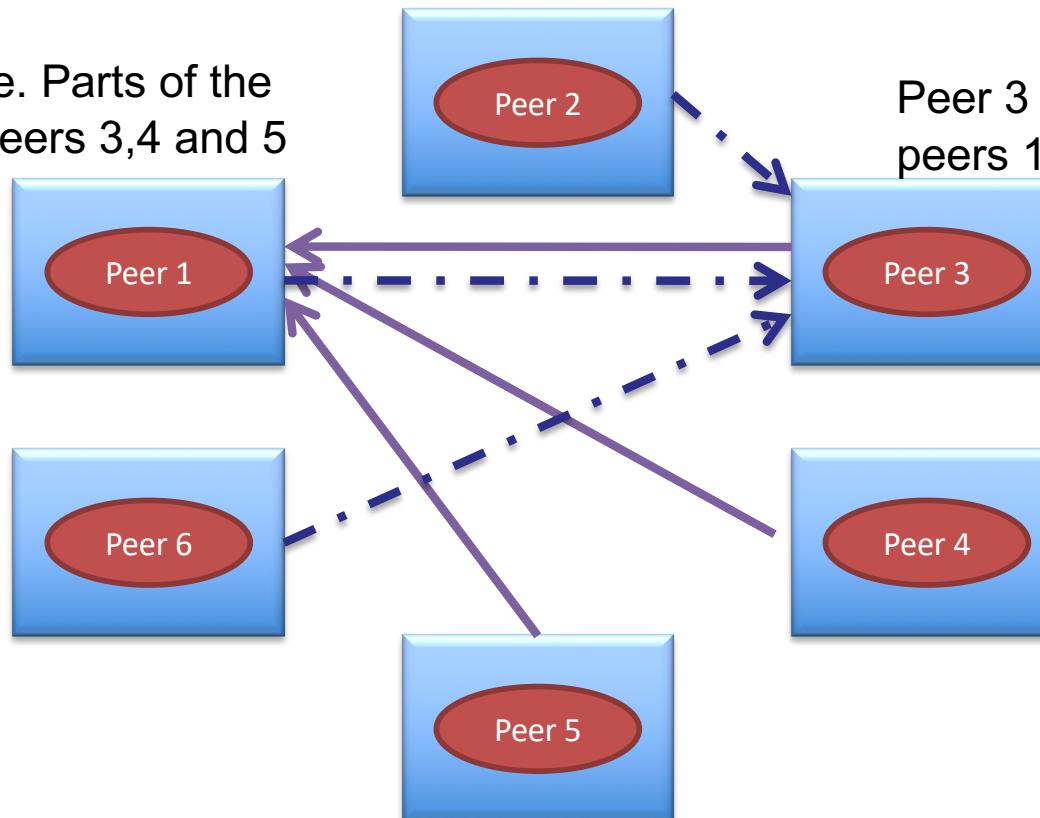


Peer to Peer Architecture

Example: Downloading files from BitTorrent

Peer 1 wants a file. Parts of the file is present at peers 3,4 and 5

Peer 3 wants a file stored at peers 1,2 and 6

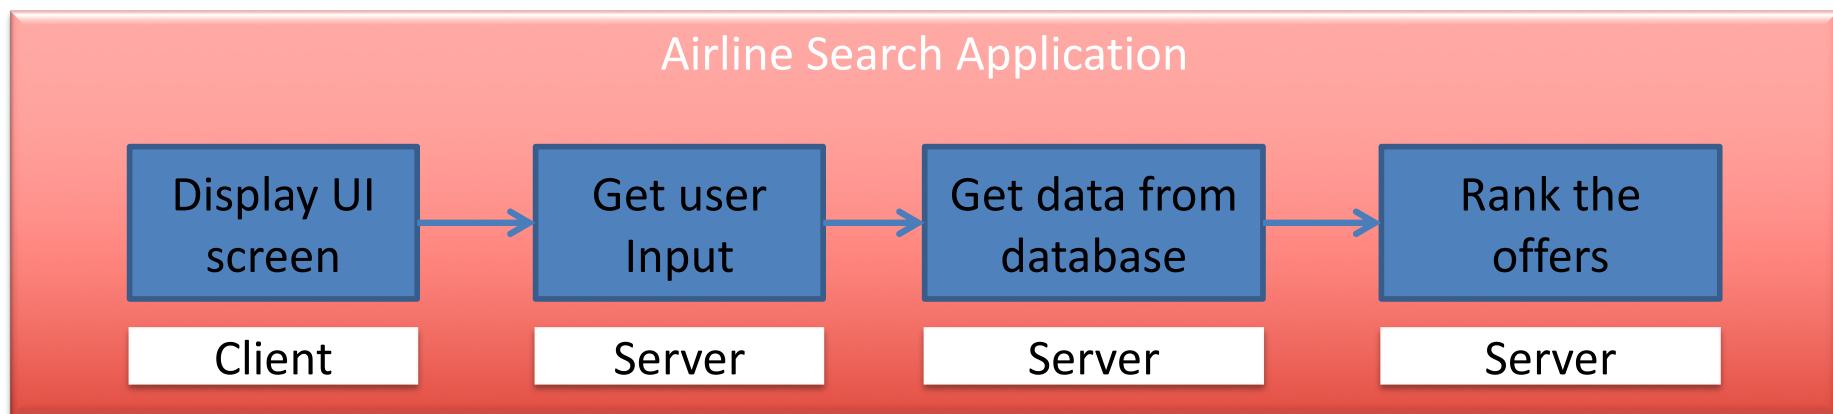


Architectural Patterns

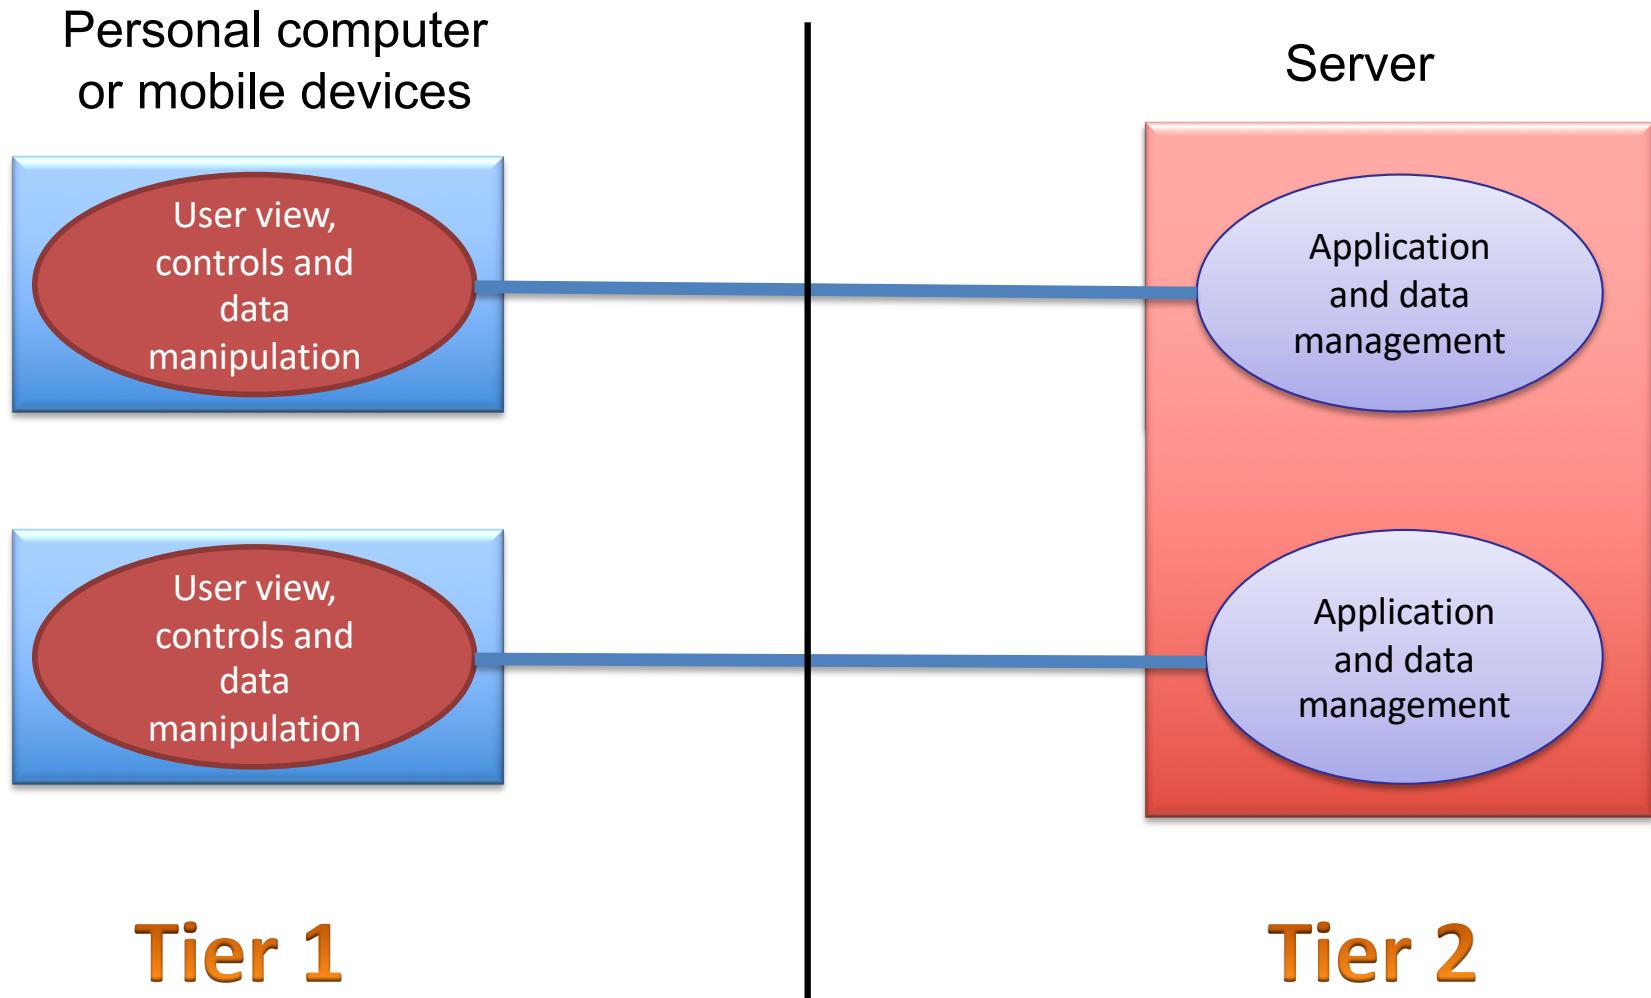
- Primitive architectural elements can be combined to form various patterns
 - Tiered Architecture
 - Layering
- Tiered architecture and layering are complementary
 - Layering = horizontal splitting of services
 - Tiered Architecture = vertical splitting of services

Tiered Architecture

- A technique to:
 1. Organize the functionality of a service, and
 2. Place the functionality into appropriate servers

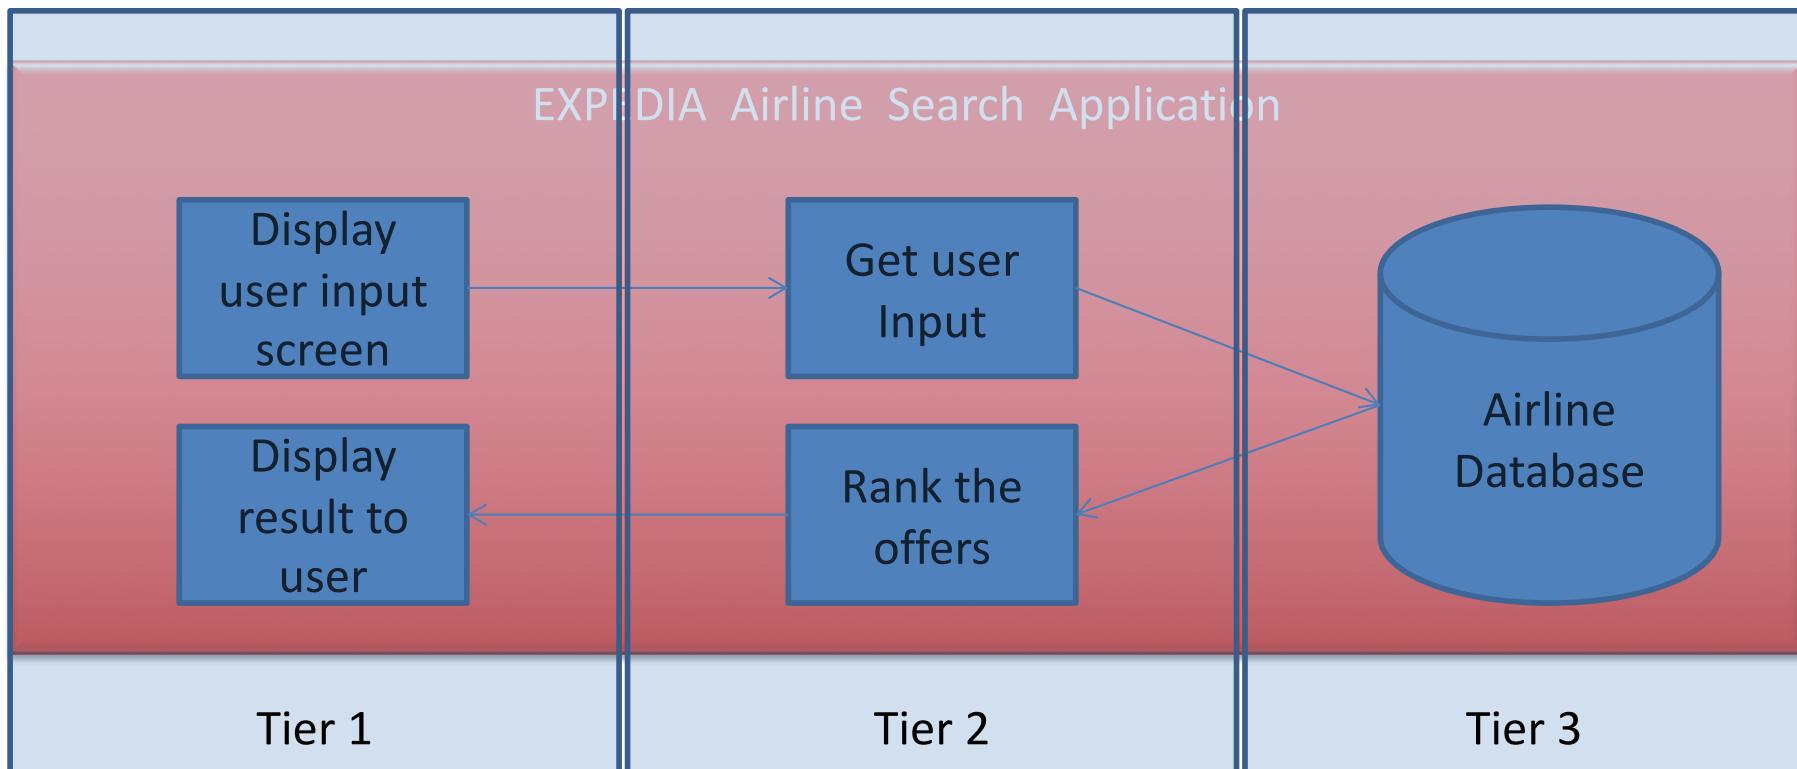


A Two-Tiered Architecture



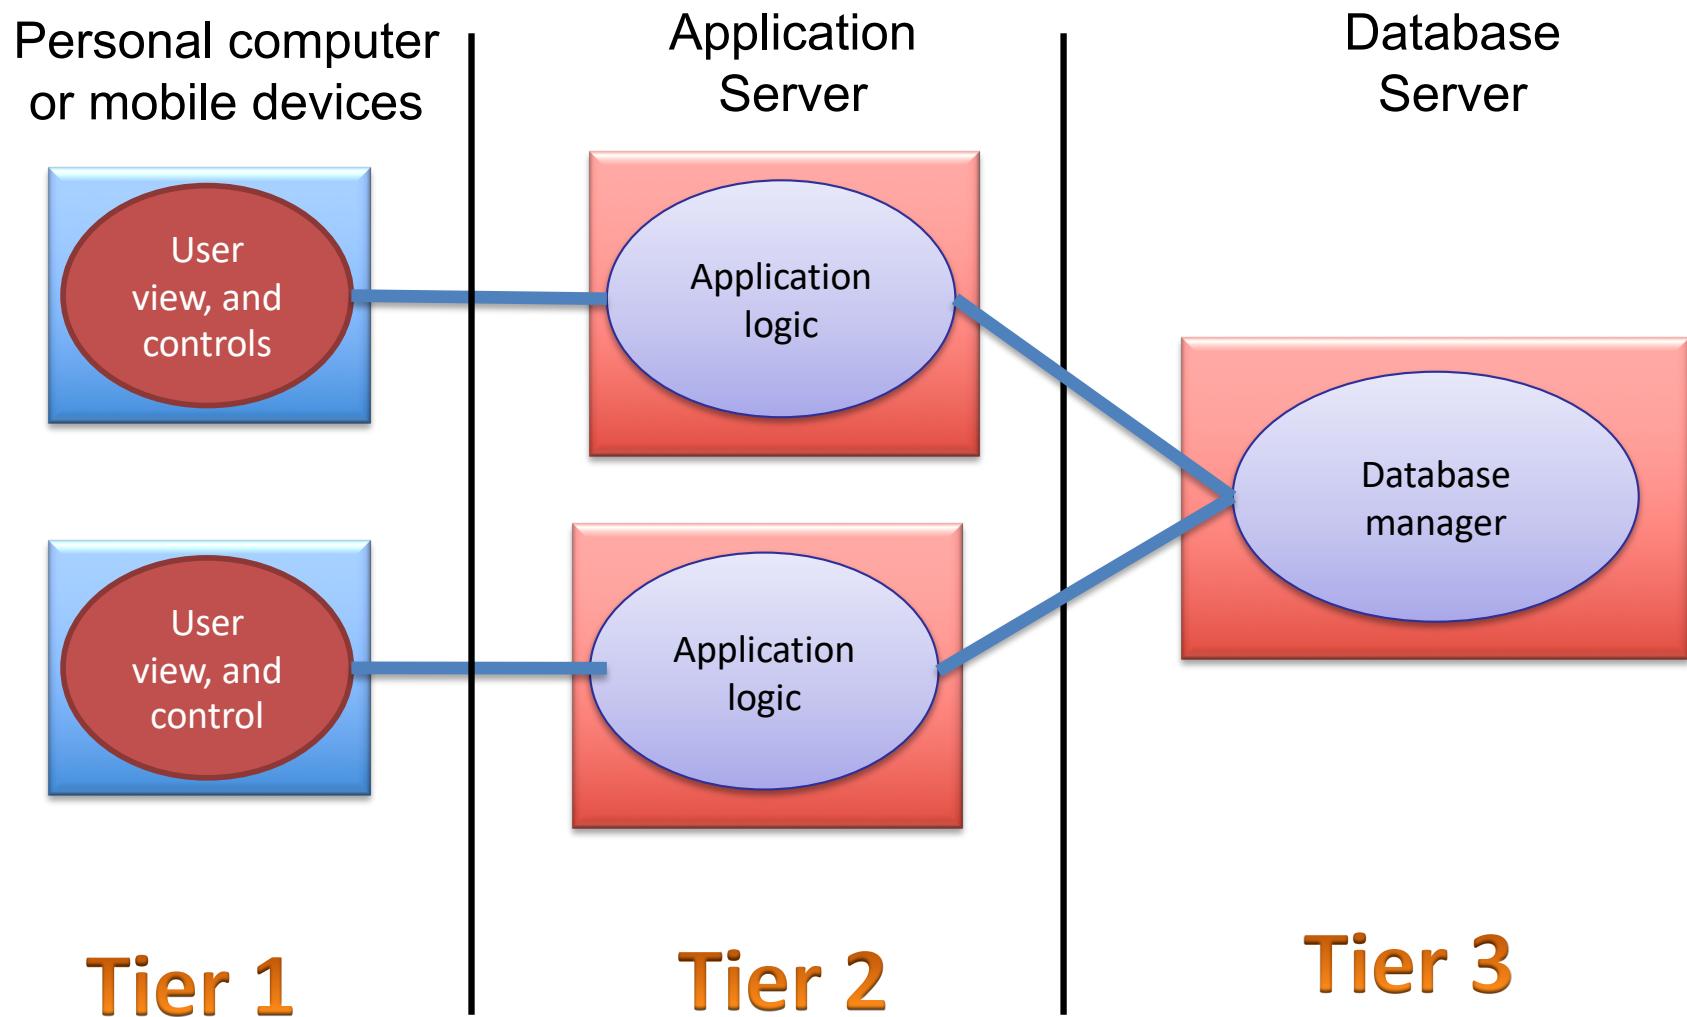
A Three-Tiered Architecture

- How do you design an airline search application:



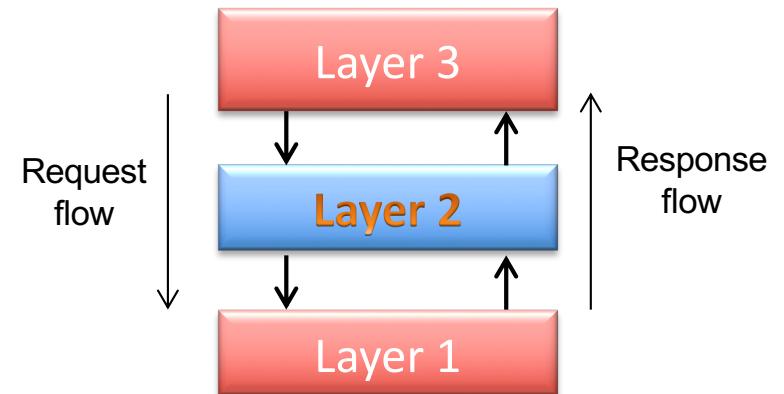
- Organize functionality of a given layer

A Three-Tiered Architecture



Layering

- A complex system is partitioned into layers
 - Upper layer utilizes the services of the lower layer
 - A horizontal organization of services
- Layering simplifies design of complex distributed systems by hiding the complexity of below layers
- Control flows from layer to layer



Layering – Platform and middleware

Distributed Systems can be organized into three layers

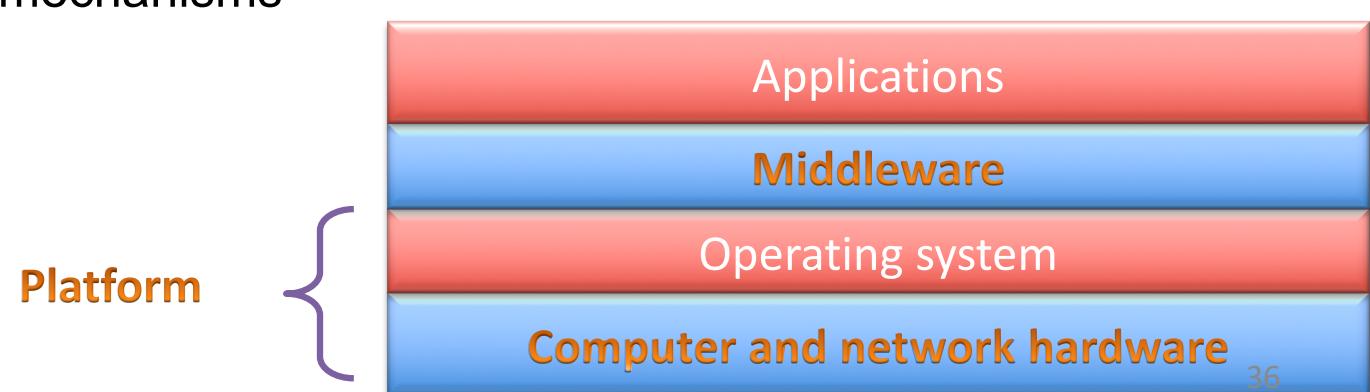
1. Platform

- Low-level hardware and software layers
- Provides common services for higher layers

2. Middleware

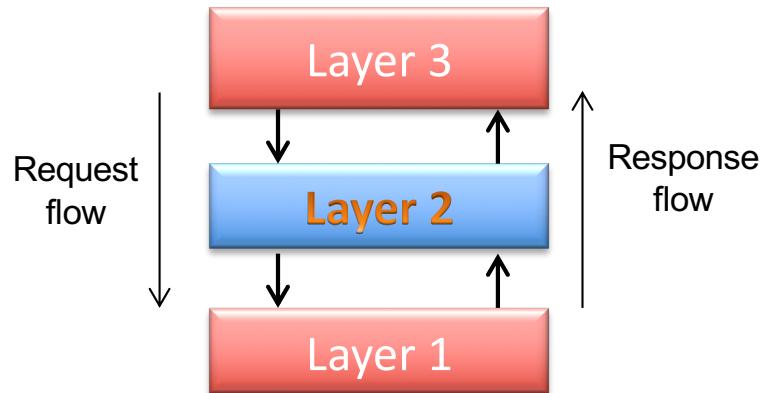
- Mask heterogeneity and provide convenient programming models to application programmers
- Typically, it simplifies application programming by abstracting communication mechanisms

3. Applications

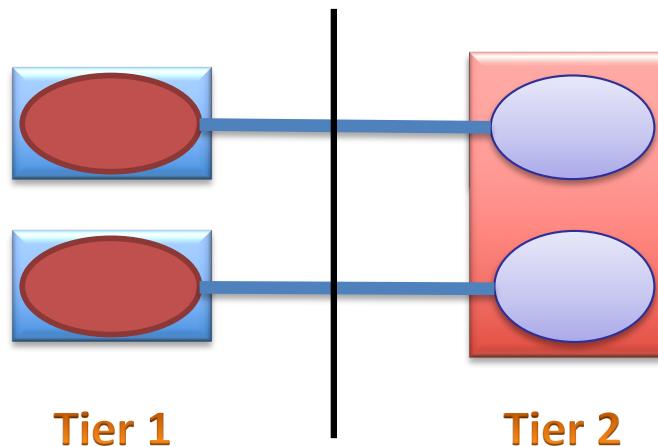


Layering vs. Tiers

- Layers: Horizontal



- Tiers: Vertical



Most large-scale DS utilize
a combination of both

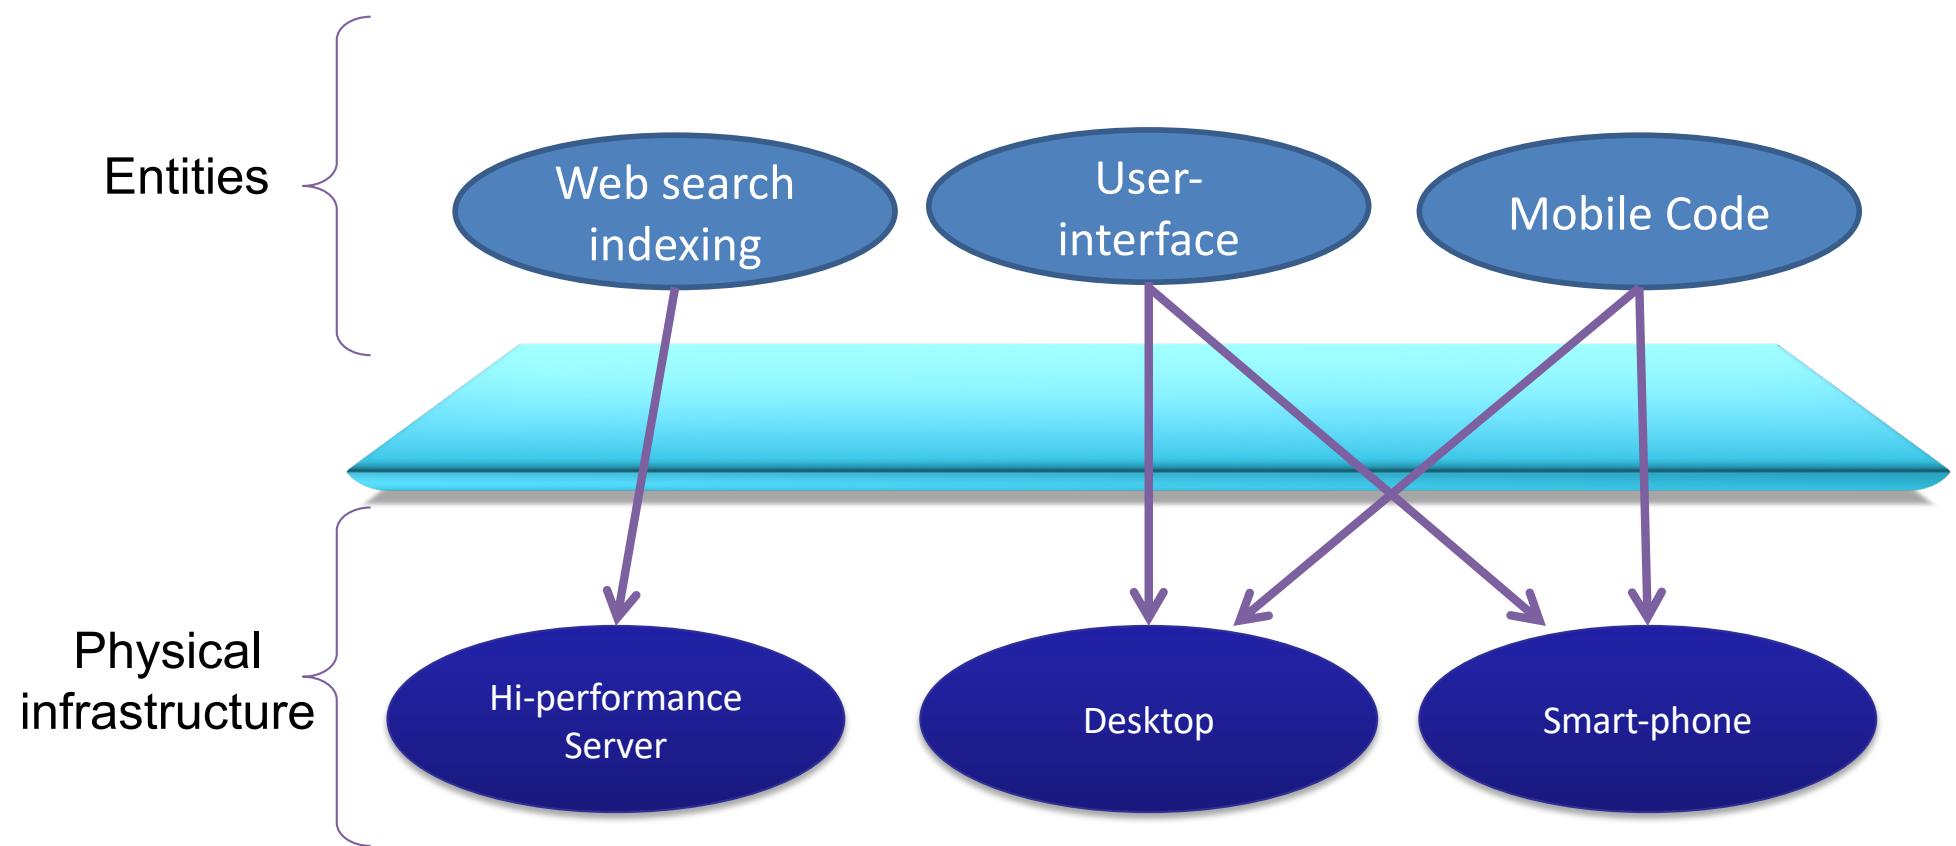
Classification of Distributed Systems

- What are the entities that are communicating in a DS?
 - a) Communicating entities
- How do the entities communicate?
 - b) Communication paradigms
- What roles and responsibilities do they have?
 - c) Roles and responsibilities
- How are they mapped to the physical distributed infrastructure?
 - d) **Placement of entities**

Placement

- Observation:
 - A large number of heterogeneous hardware (machines, network).
 - Smart mapping of entities (processes, objects) to hardware helps performance, security and fault-tolerance.
- “Placement” maps entities to underlying physical distributed infrastructure.
 - Placement should be decided after a careful study of application characteristics
 - Example strategies:
 - Mapping services to multiple servers
 - Moving the mobile code to the client

Placement



Placement

- Why are there data centers all over the planet
 - Placement!
 - Goal:
 - Place data and service close to the user
 - Better performance (delay, bandwidth, ...)
- Example
 - Videos that go “viral”
 - Where to place them?
 - Data travels with the daytime
 - From Europe to the US to Asia to ...

Recap Architecture

- So far, we have covered primitive architectural elements
 - Communicating entities
 - Communication paradigms of entities
 - [Middleware, Indirect Communication](#)
 - Roles and responsibilities that entities assume, and resulting architectures
 - [Client-Server, Peer-to-Peer, Hybrid](#)
 - Placement of entities

Part II

Processes & Concurrency

Concurrency

Example

```
int y = 0, z = 0;  
  
thread A {  
    x = y + z;  
}  
  
thread B {  
    y = 1; z = 2;  
}
```

- OS Scheduler
 - can interrupt a thread at any point in time
 - and schedule another thread

Concurrency

- Possible results for x?

Example

```
int y = 0, z = 0;
```

```
thread A {  
    x = y + z;  
}
```

```
thread B {  
    y = 1; z = 2;  
}
```



Concurrency

Example

```
int y = 0, z = 0;
```

```
thread A {  
    x = y + z;  
}
```

```
thread B {  
    y = 1; z = 2;  
}
```

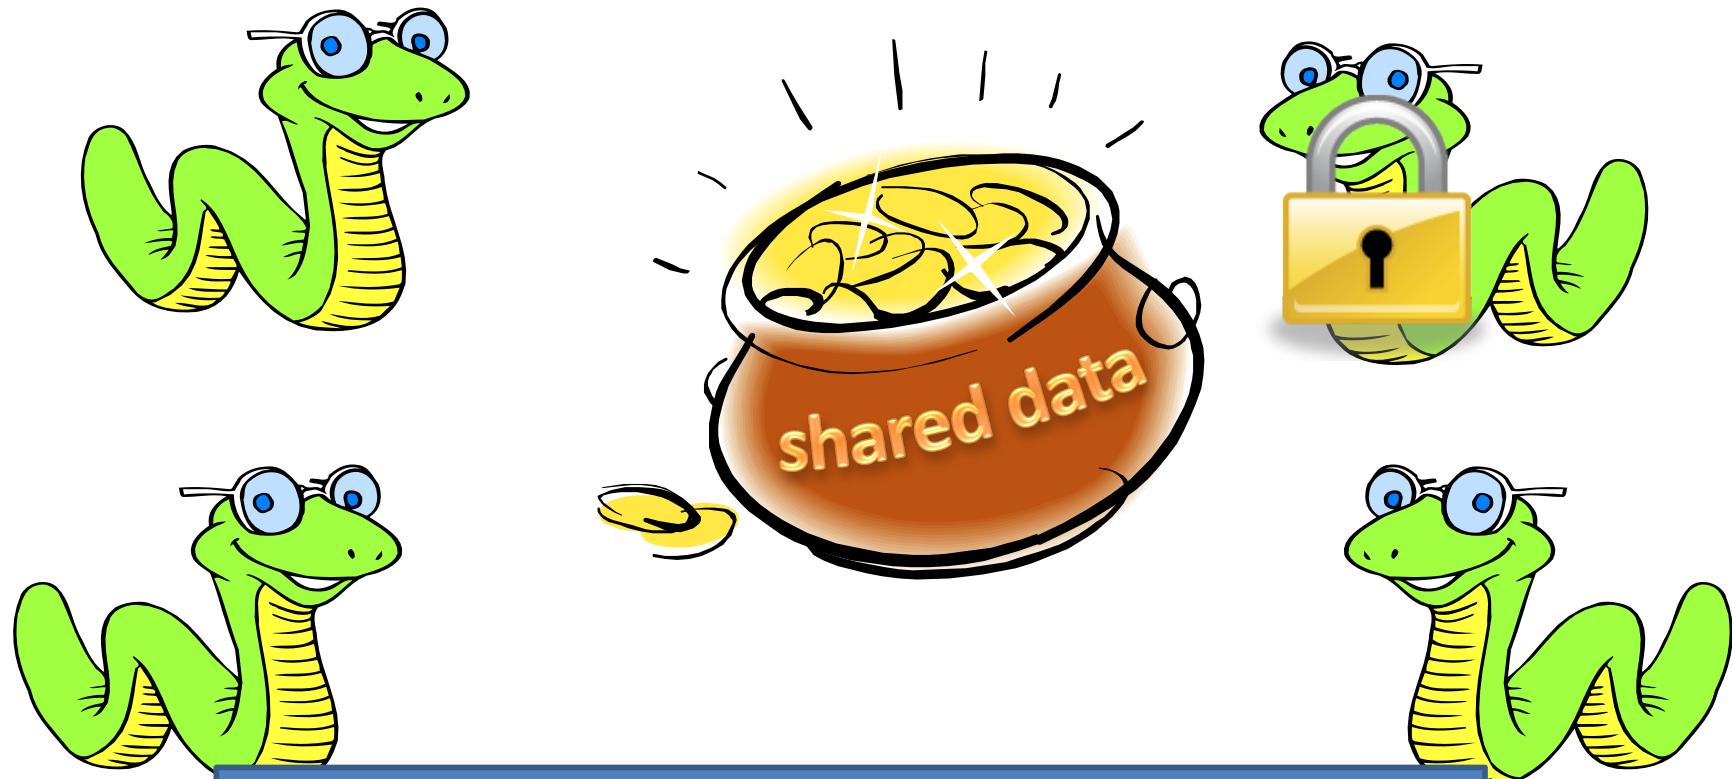
- Possible results for x?
 - 0, 1, 2, 3
- What causes this problem?
 - Concurrency
- How to fix?
 - Use locks

Thread Synchronization



If lock is not available, threads wait
-> Execution becomes serialized (one after another gets access)

Thread Synchronization



What can go wrong? Many things..

- We forget to release a lock (bug)
- What happens when the process holding a lock crashes?
-

But

- In distributed settings
 - Multiple machine or processes

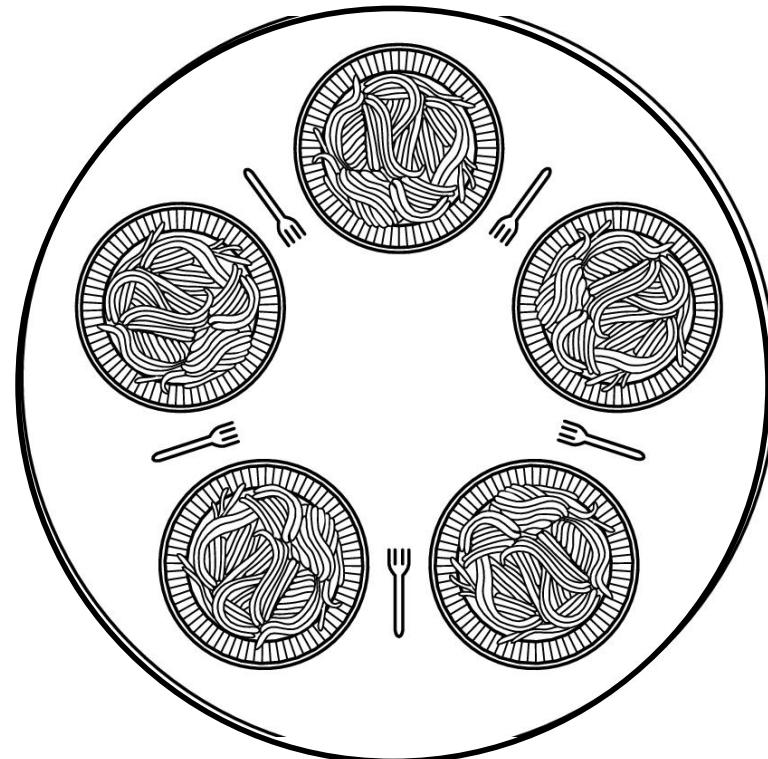
Locks alone are not always sufficient!

Dining Philosophers

- A problem that was invented to illustrate a access to shared resources
- Our focus here is on the notion of sharing resources that *only one user at a time* can own
 - Such as a keyboard on a machine with many processes active at the same time
 - Or a special disk file that only one can write at a time (bounded buffer is an instance)

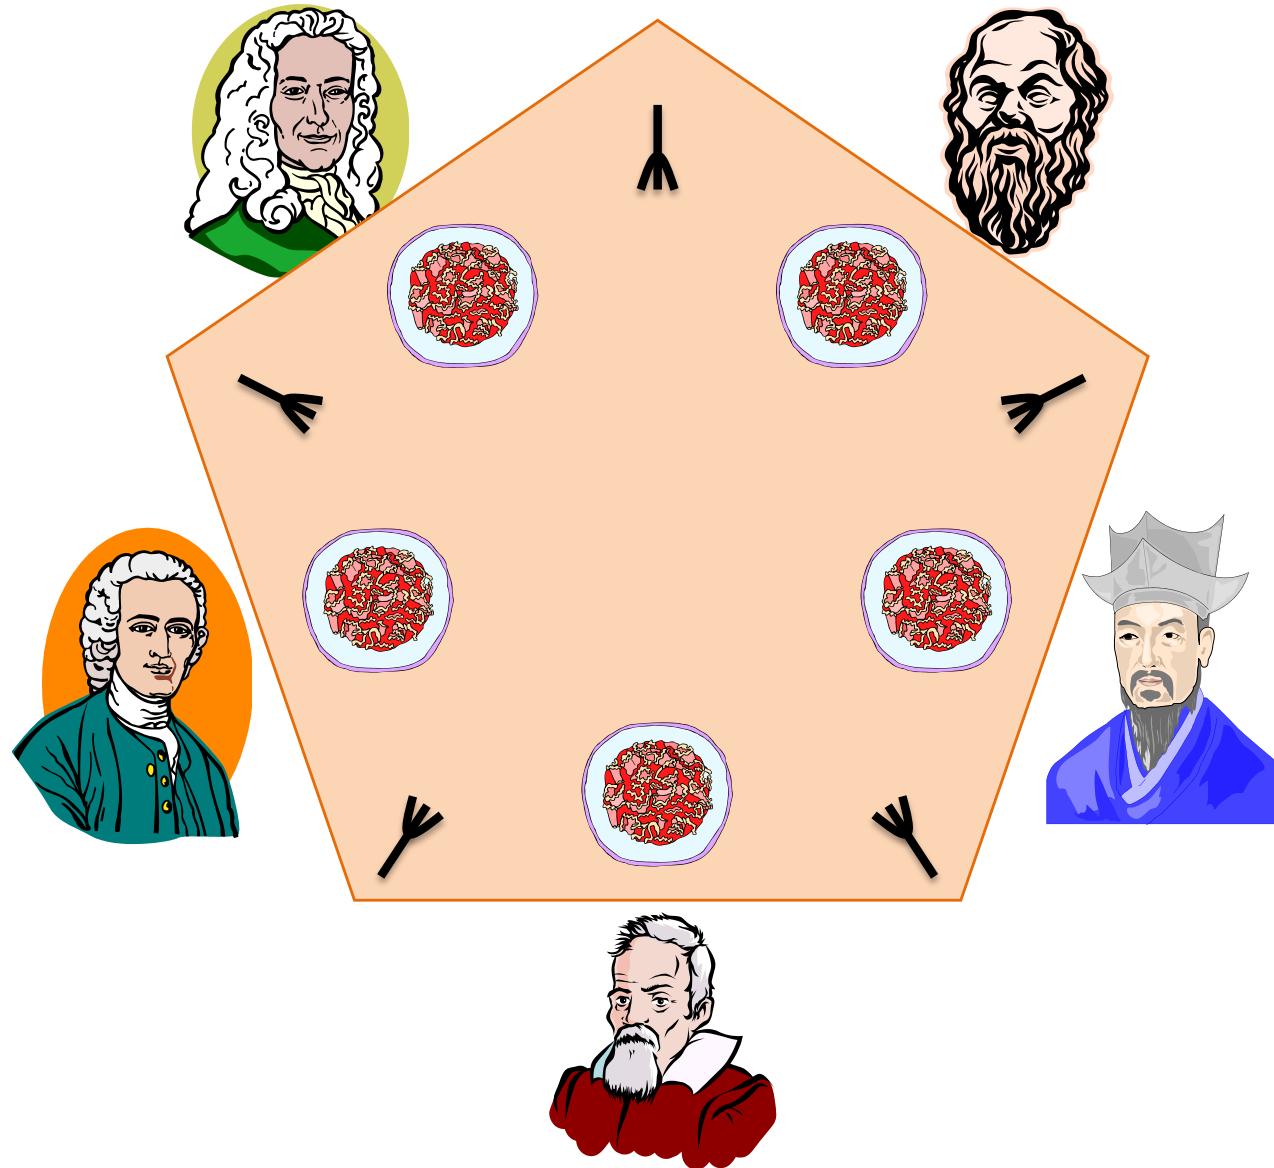
Dining Philosopher's Problem

- Invented by Dijkstra in 1965
 - As exam question ;-)
- Philosophers eat/think
- N Philosophers, N forks
- Eating needs two forks
- Pick one fork at a time



Idea is to capture the concept of multiple processes competing for limited resources

Dining Philosopher's Problem



Ok!

Not good!

Rules of the Game

- The philosophers are very logical
 - They want to settle on a shared policy that all can apply concurrently
 - They are hungry: the policy should let everyone eat (eventually)
 - They are utterly dedicated to the proposition of equality: the policy should be totally fair

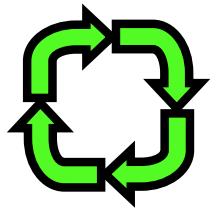
Note

- We want fully distributed solution
 - Do not want philosophers to make arrangements with their neighbors
 - Cascade: these would also need to make arrangements with their neighbors
 - This would limit scalability and make things complex

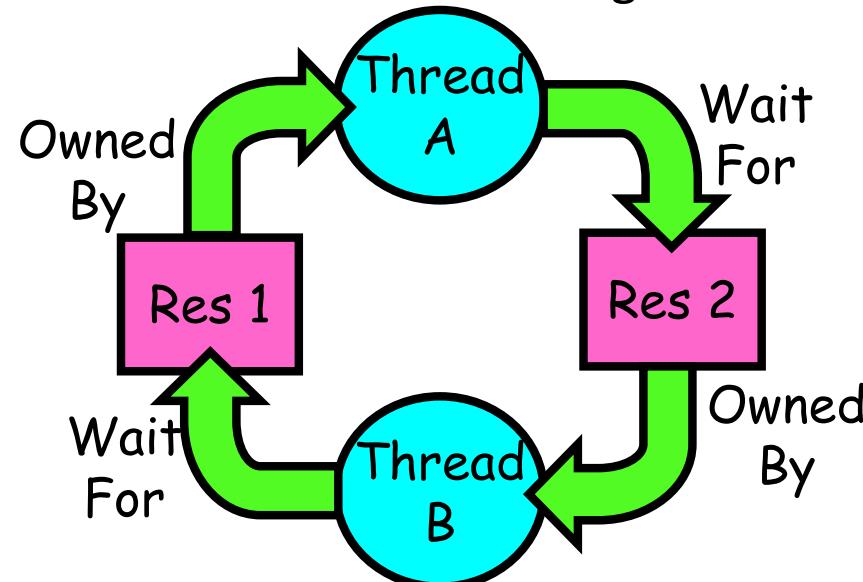
What can go wrong?

- Primarily, we worry about?
 - Starvation: A policy that can leave some philosopher hungry in some situation (even one where the others collaborate)
 - Deadlock: A policy that leaves all the philosophers “stuck”, so that nobody can do anything at all
 - Livelock: A policy that makes them all do something endlessly without ever eating!

Starvation vs Deadlock



- Starvation vs. Deadlock
 - Starvation: thread waits indefinitely
 - Example, low-priority thread waiting for resources constantly in use by high-priority threads
 - Deadlock: circular waiting for resources
 - Thread A owns Res 1 and is waiting for Res 2
Thread B owns Res 2 and is waiting for Res 1



- Deadlock \Rightarrow Starvation but not vice versa
 - Starvation can end (but doesn't have to)
 - Deadlock can't end without external intervention

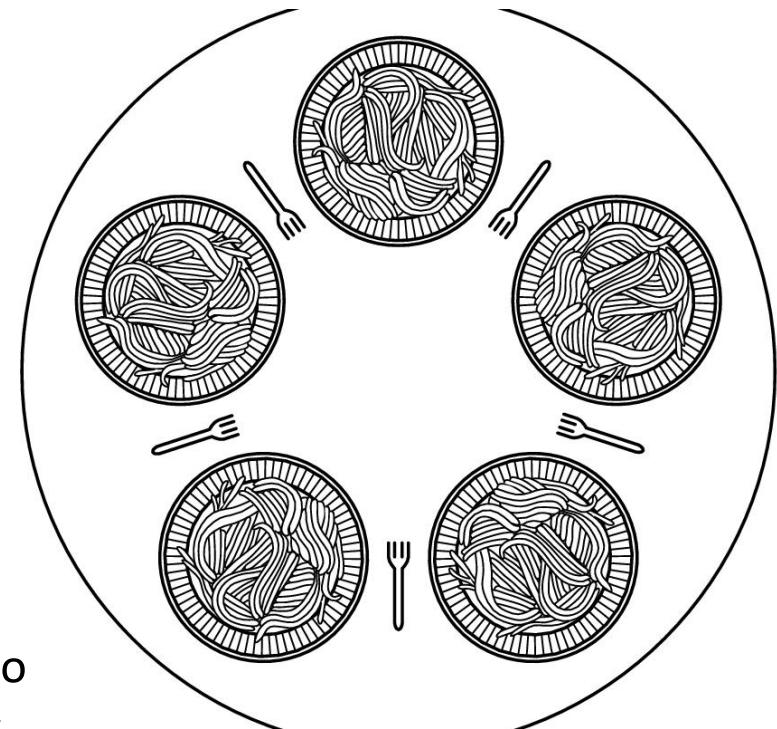
A flawed conceptual solution

```
# define N      5
```

Philosopher i (0, 1, .. 4)

```
do {  
    think();  
    take_fork(i);  
    take_fork((i+1)%N);  
    eat(); /* yummy */  
    put_fork(i);  
    put_fork((i+1)%N);  
} while (true);
```

Oops! Subject to deadlock if they all pick up their “right” fork simultaneously!

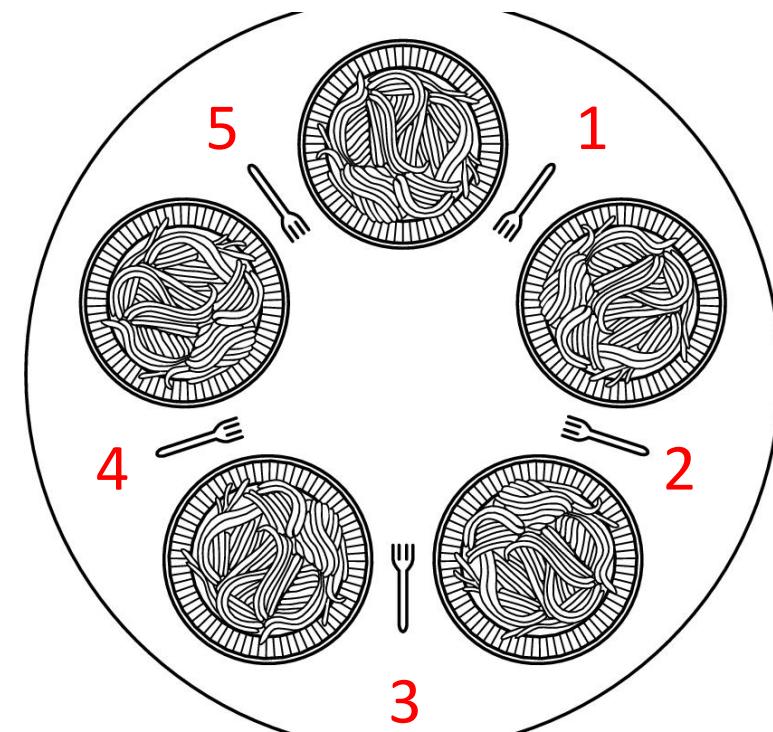


Solutions?

- Central entity: waiter
 - Assigns who can eat
 - Problem?
 - Does not scale, single point of failure
- Distributed
 - Naïve “Solution”:
 - Talking to determine who can eat
 - Requires global state (does not scale)
 - Solution 1: Assign order to resources (see next slides)
 - Solution 2: Probabilistic (see next slides)
 - Add randomization

Solution 1: Partial Order

- Assign order to forks: 1 to n
 - Always try to pick the resource with the lower ID first
 - Avoids deadlocks, starvation, ...
 - But...
 - Not always practicable to number resources



Solution 2: Probabilistic Solution

- Assume Random “coin toss”
 - Guaranteed with probability 1 to break symmetry
- Idea: Try to get one first
 - Then get other
 - If can’t get other, put first down and try again
- But don’t go for the same fork first every time

Think

trying = true

While *trying*

s = random(left,right)

Wait for fork *s* then take it

If fork $\sim s$ available

take it

trying = false

else

drop fork *s*

Eat

drop *s*=random(left,right)

drop fork $\sim s$

Lehman, D. and Michael O. Rabin. "On the advantages of free choice: A symmetric and fully distributed solution of the dining philosophers problem," Proceedings of the Eighth ACM Symposium on the Principles of Programming Languages. (1981) 60

Recap Concurrency

- Concurrency
 - Once multiple entities operate on the same data
 - Processes, multiple machines, ...
 - Use Locks
 - Avoid deadlocks, starvation, ...
- Concurrency is natural in distributed systems
 - See your labs ;-)
 - Important to understand concurrency on a single system
 - Before moving on to large-scale DS

Next Time

- Naming

Questions?