

Research Group
Distributed Systems

C | A | U

Christian-Albrechts-Universität zu Kiel

Technische Fakultät

Distributed Systems

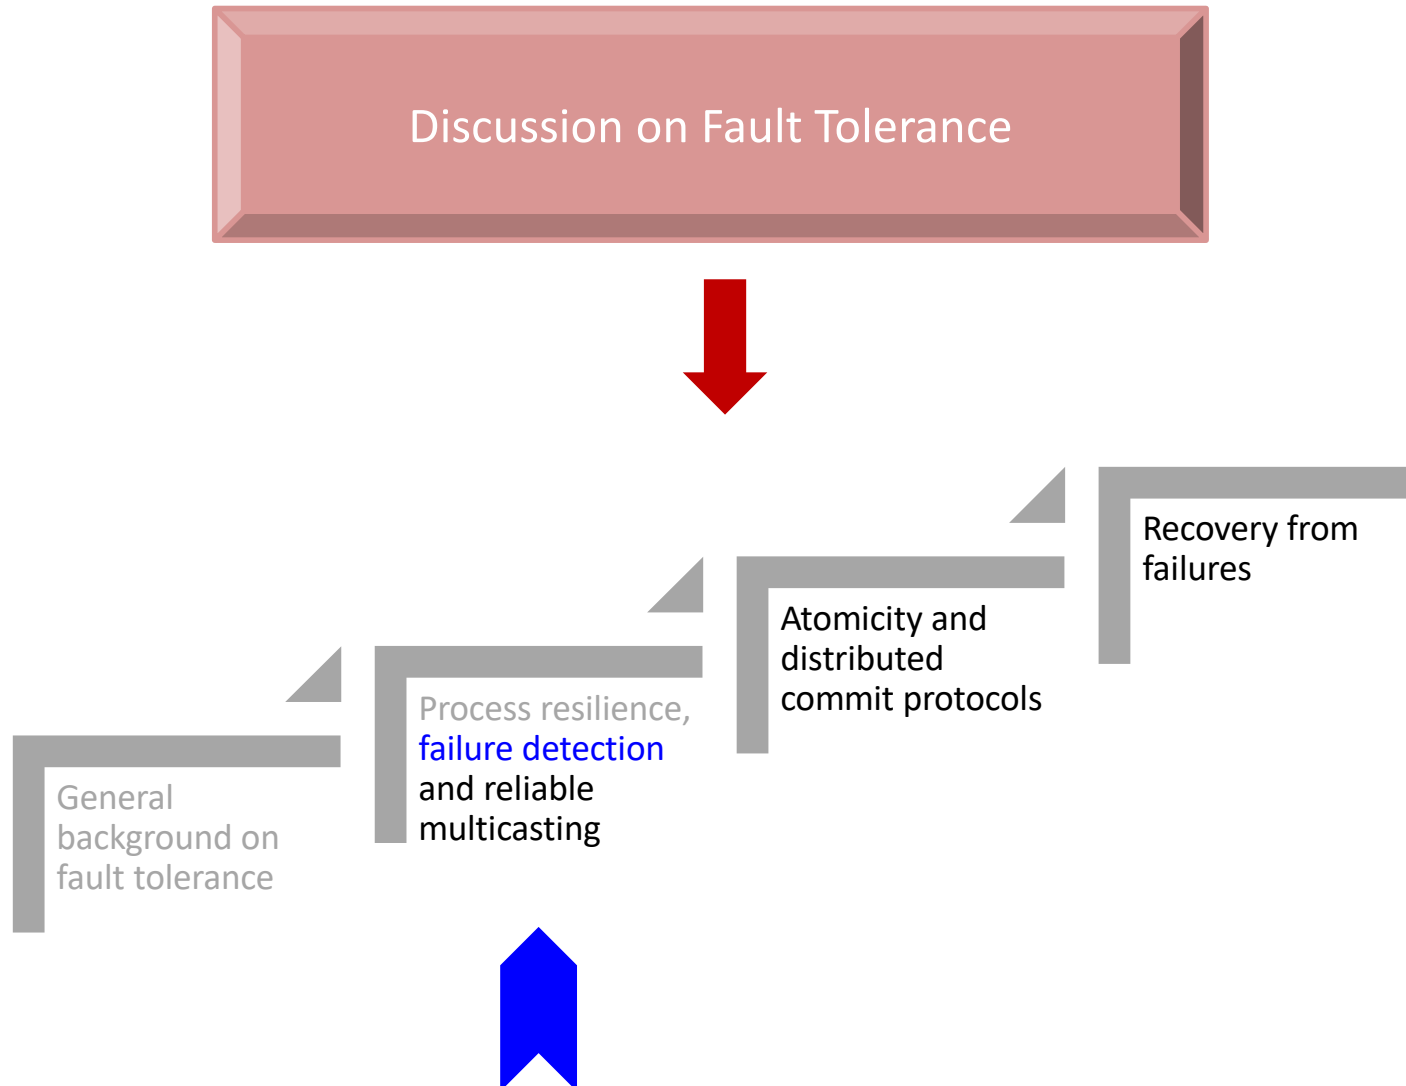
Fault-Tolerance II

Olaf Landsiedel

Last Time

- Fault Tolerance I
 - Failures
 - Process Resilience:
 - Byzantine Generals

Objectives



Process Failure Detection


- Before we properly mask failures, we generally need to detect them
- For a group of processes, non-faulty members should be able to decide who is still a member and who is not
- Two policies:
 - Processes actively send “are you alive?” messages to each other (i.e., *pinging each other*)
 - Processes passively wait until messages come in from different processes

Timeout Mechanism

- In failure detection a *timeout mechanism* is usually involved
 - Specify a timer, after a period of time, trigger a timeout
 - Best known example?
 - TCP
 - Problem?
 - However, due to unreliable networks, simply stating that a process has failed because it does not return an answer to a ping message may be wrong

Example: FUSE

- In FUSE, processes can be joined in a group that spans a WAN
 - If one fails, others shall switch to fail state, too
- The group members create a spanning tree that is used for monitoring member failures
- An active (pinging) policy is used where a single node failure is rapidly promoted to a group failure notification



● Failed Member ● Alive Member

Failure Considerations

- There are various issues that need to be taken into account when designing a failure detection subsystem:
 1. Failure detection can be done as a side-effect of regularly exchanging information with neighbors (e.g., **gossip-based information dissemination**)
 2. A failure detection subsystem should ideally be able to distinguish network failures from node failures
 3. When a member failure is detected, how should other non-faulty processes be informed


Summary

- Always design for failures
 - See your labs, see data centers
- Often the reason for failure is different than expected
 - See data centers
- Algorithms that can deal well with failures
 - Are commonly quite complex

Leslie Lamport

“A distributed system is one in which the failure of a machine you have never heard of can cause your own machine to become unusable”

- Issue is dependency on critical components
- Notion is that state and “health” of system at site A is linked to state and health at site B

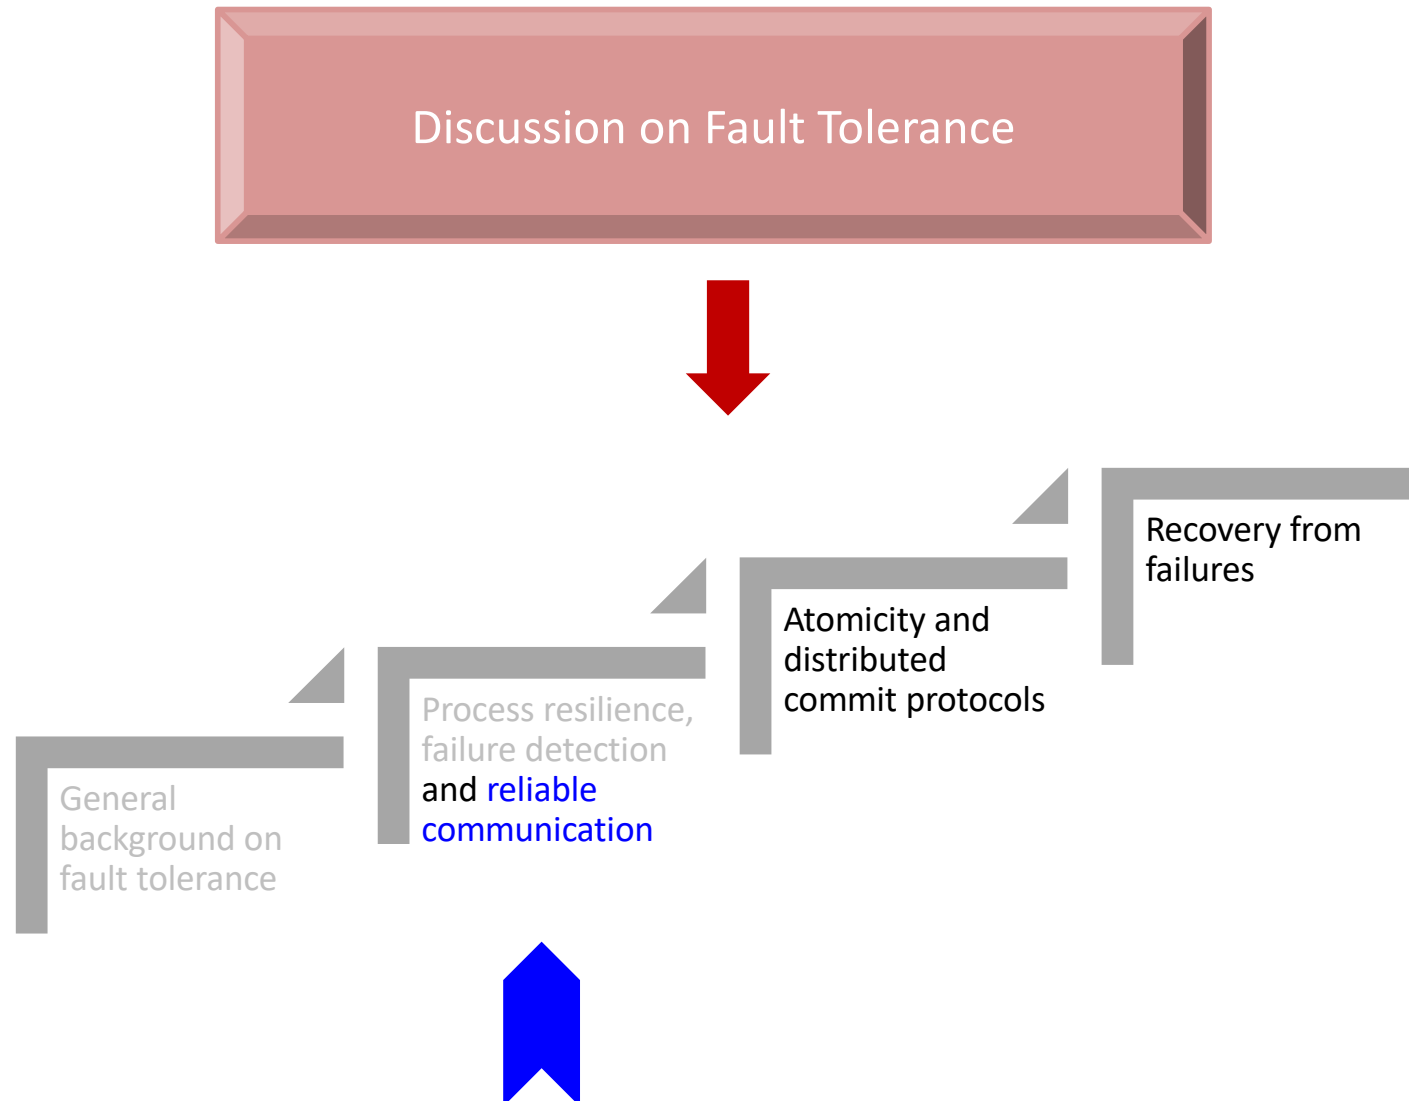


Picture: <http://research.microsoft.com/en-us/um/people/lamport/>

Today...


- Today's session
 - Fault Tolerance – *Part II*
 - Reliable communication

Objectives



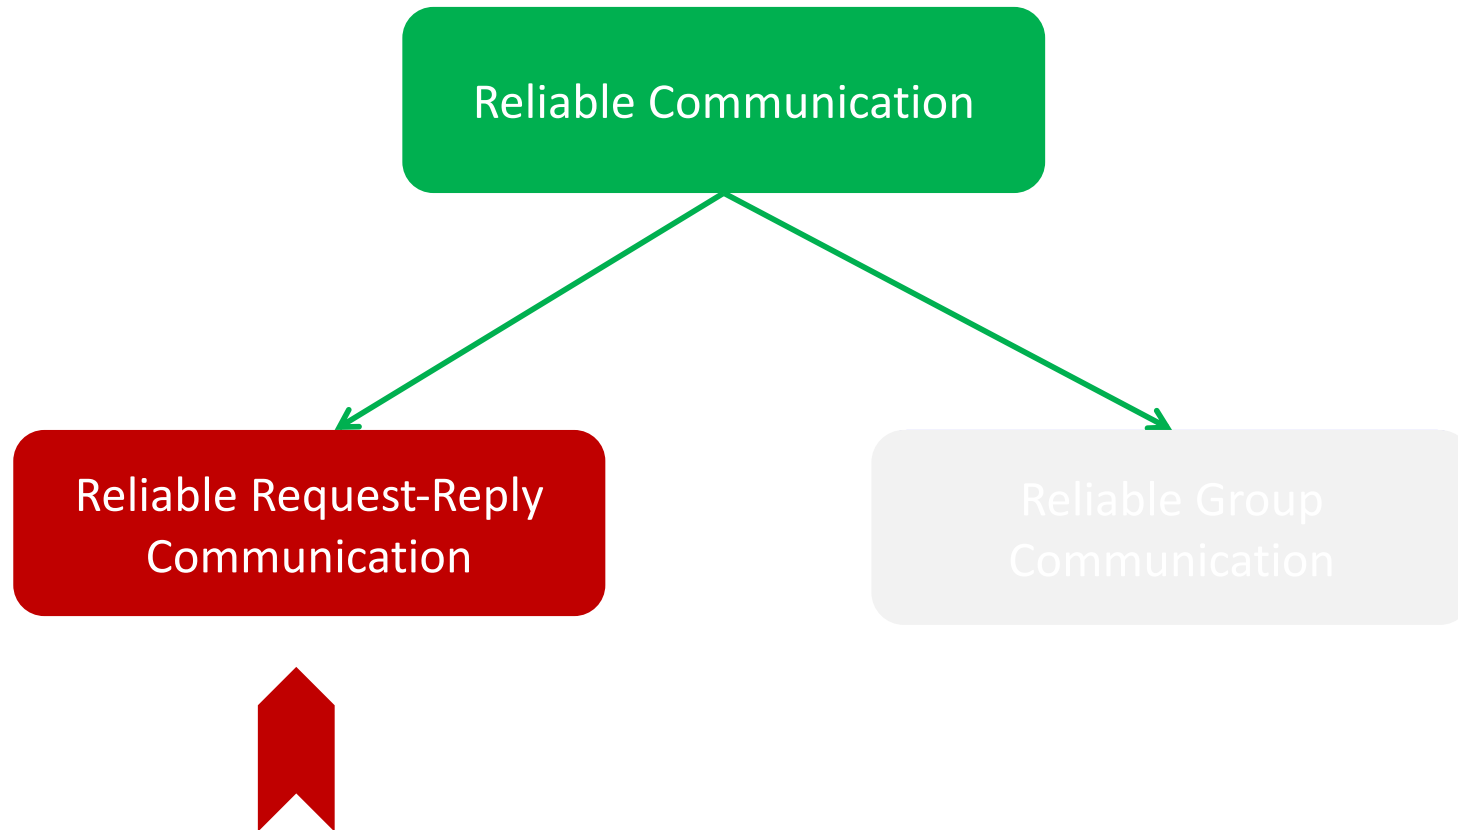
Reliable Communication

- Fault tolerance in distributed systems typically concentrates on faulty processes



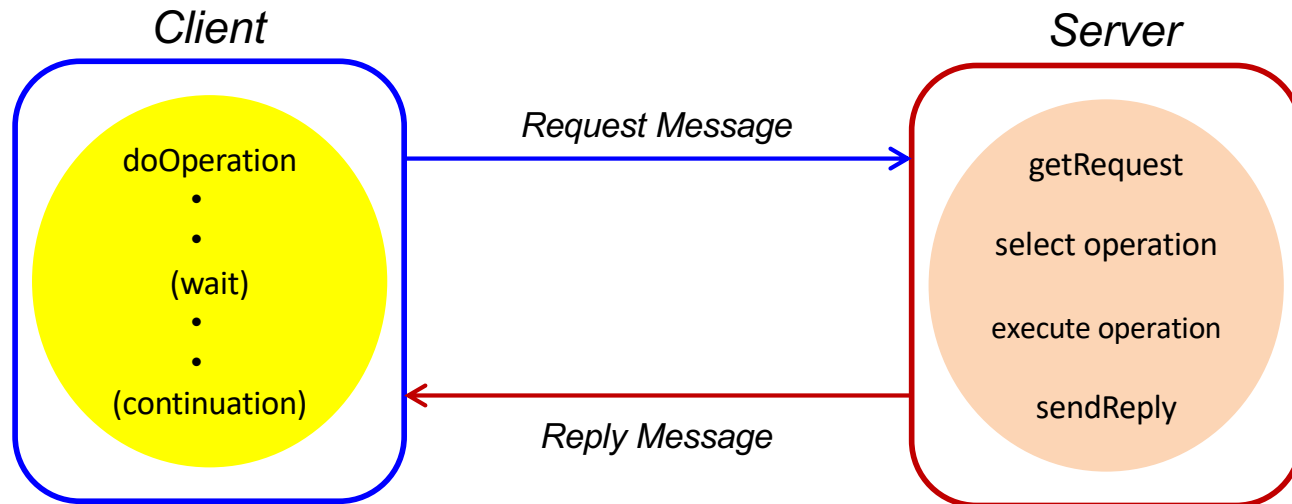
- However, we also need to consider communication failures
- We will focus on two types of reliable communication:
 - Reliable request-reply communication (e.g., RPC)
 - Reliable group communication (e.g., multicasting schemes)

Reliable Communication



Request-Reply Communication

- The request-reply communication is designed to support the roles and message exchanges in typical client-server interactions



- This sort of communication is mainly based on a trio of communication primitives, *doOperation*, *getRequest* and *sendReply*

Timeouts

- Request-reply communication may suffer from *crash*, *omission*, *timing*, and *byzantine* failures
- To allow for occasions where a request or a reply message is not delivered (e.g., lost), *doOperation* uses a *timeout*
- There are various options as to what *doOperation* can do after a timeout:
 - Return immediately with an indication to the client that the request has failed
 - Send the request message repeatedly until either a reply is received or the server is assumed to have failed

Duplicate Filtering

- In cases when the request message is retransmitted, the server may receive it more than once
- This can cause the server executing an operation more than once for the same request
- To avoid this, the server should recognize successive messages from the same client and filter out duplicates

Lost Reply Messages

- If the server has already sent the reply when it receives a duplicate request, it can either:
 - Re-execute the operation again to obtain the result
 - Or do not re-execute the operation if it has chosen to retain the outcome of the first execution
- Not every operation can be executed more than once and obtain the same results each time
- An idempotent operation is an operation that can be performed repeatedly with the same effect as if it had been performed exactly once

History

- For servers that require retransmission of replies without re-execution of operations, a history may be used
- The term 'history' is used to refer to a structure that contains a record of (reply) messages that have been transmitted

Fields of a history record:

Request ID

Message

Client ID

History

- The server can interpret each request from a client as an ACK of its previous reply
 - Thus, the history needs contain only the last reply message sent to each client
- But, if the number of clients is large, memory cost might become a problem
- Messages in a history are normally discarded after a limited period of time

Summary of Request-Reply Protocols

- The *doOperation* can be implemented in different ways to provide different delivery guarantees. The main choices are:
 1. Retry request message: Controls whether to retransmit the request message until either a reply is received or the server is assumed to have failed
 2. Duplicate filtering: Controls when retransmissions are used and whether to filter out duplicate requests at the server
 3. Retransmission of results: Controls whether to keep a history of result messages to enable lost results to be retransmitted without re-executing the operations at the server

Request-Reply Call Semantics

- Combinations of request-reply protocols lead to a variety of possible semantics for the reliability of remote invocations

Fault Tolerance Measure			Call Semantics
Retransmit Request Message	Duplicate Filtering	Re-execute Procedure or Retransmit Reply	
No	N/A	N/A	<i>Maybe</i>
Yes	No	Re-execute Procedure	<i>At-least-once</i>
Yes	Yes	Retransmit Reply	<i>At-most-once</i>

Maybe Semantics

- With *maybe semantics*, the remote procedure call may be executed once or not at all
- *Maybe semantics* arises when no fault-tolerance measures are applied and can suffer from the following types of failure:
 - Omission failures if the request or result message is lost
 - Crash failures when the server containing the remote operation fails
- *Maybe semantics* is useful only for applications in which occasional failed calls are acceptable

Maybe Semantics: Revisit

Fault Tolerance Measure			Call Semantics
Retransmit Request Message	Duplicate Filtering	Re-execute Procedure or Retransmit Reply	
No	N/A	N/A	<i>Maybe</i>
Yes	No	Re-execute Procedure	<i>At-least-once</i>
Yes	Yes	Retransmit Reply	<i>At-most-once</i>

At-Least-Once Semantics

- With at-least-once semantics, the invoker keeps retransmitting the request message until a reply is received
- At-least-once semantics:
 - Masks the omission failures due to retransmissions
 - Suffers from crash failures when the server containing the remote operation fails
 - Might suffer from response failures if a re-executed operation is not idempotent

At-Least-Once Semantics: Revisit

Fault Tolerance Measure			Call Semantics
Retransmit Request Message	Duplicate Filtering	Re-execute Procedure or Retransmit Reply	
No	N/A	N/A	<i>Maybe</i>
Yes	No	Re-execute Procedure	<i>At-least-once</i>
Yes	Yes	Retransmit Reply	<i>At-most-once</i>

At-Most-Once Semantics

- With at-most-once semantics, the invoker *gives up immediately and reports back a failure*
- At-most-once semantics:
 - Masks the omission failures of the request or result messages by retransmitting request messages
 - Avoids response failures by ensuring that each operation is never executed more than once

At-Most-Once Semantics: Revisit

Fault Tolerance Measure			Call Semantics
Retransmit Request Message	Duplicate Filtering	Re-execute Procedure or Retransmit Reply	
No	N/A	N/A	<i>Maybe</i>
Yes	No	Re-execute Procedure	<i>At-least-once</i>
Yes	Yes	Retransmit Reply	<i>At-most-once</i>

Classes of Failures in Request-Reply Communication

- There are 5 different classes of failures that can occur in request-reply systems:
 1. The client is unable to locate the server
 2. The request message from the client to the server is lost
 3. The server crashes after receiving a request
 4. The reply message from the server to the client is lost
 5. The client crashes after sending a request

Classes of Failures in Request-Reply Communication

- There are 5 different classes of failures that can occur in request-reply systems:
 1. The client is unable to locate the server
 2. The request message from the client to the server is lost
 3. The server crashes after receiving a request
 4. The reply message from the server to the client is lost
 5. The client crashes after sending a request

Possible Solution

- One possible solution for the client being unable to locate the server is to have *doOperation* raise an exception at the client side
- Considerations:
 - Not every language has exceptions or signals
 - Writing an exception identifies the location of the error and hence destroys the transparency of the distributed system

Classes of Failures in Request-Reply Communication

- There are 5 different classes of failures that can occur in request-reply systems:
 1. The client is unable to locate the server
 2. The request message from the client to the server is lost
 3. The server crashes after receiving a request
 4. The reply message from the server to the client is lost
 5. The client crashes after sending a request

Possible Solution

- The *doOperation* can start a timer when sending the request message
- If the timer expires before a reply or an ACK comes back, the message is sent again
- Considerations:
 - If the message was lost, the server might not be able to recognize the difference between a first transmission and a retransmission
 - If the message was not lost, the server has to detect that it is dealing with a retransmission request

Classes of Failures in Request-Reply Communication

- There are 5 different classes of failures that can occur in request-reply systems:
 1. The client is unable to locate the server
 2. The request message from the client to the server is lost
 3. The server crashes after receiving a request
 4. The reply message from the server to the client is lost
 5. The client crashes after sending a request

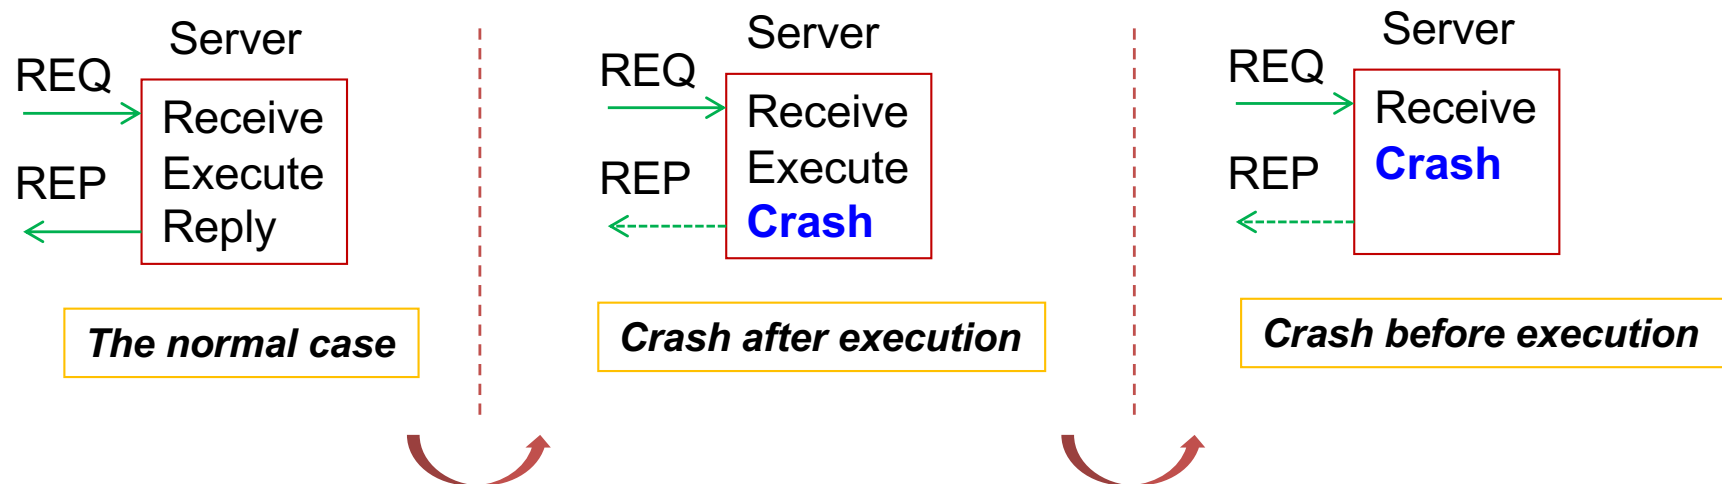
Possible Solution (1)

- The *doOperation* can start a *timer* when sending the request message
- If the timer expires before a reply or an ACK comes back, the message is sent again
- *We can apply any of the 3 request-reply call semantics*
- Considerations:
 - The crash failure may occur *either before or after* the operation at the server is executed. The *doOperation* cannot figure that out

Possible Solution (2)

- Considerations (Cont'd):

- The sequence of events at server is as follows:



- In the last 2 cases, *doOperation* cannot tell which is which. All it knows is that its *timer* has expired

A Printing Example (PE): Normal Scenario

- A client's remote operation consists of printing some text at a server
- When a client issues a request, it receives an ACK that the request has been delivered to the server
- The server sends a completion message to the client when the text is printed

PE: Possible Events at Server

- Three events can happen at the server:
 1. Send the completion message (M)
 2. Print the text (P)
 3. Crash (C)

PE: Server Strategies

- The server has a choice between two strategies:
 1. Send a completion message just before commanding the printer to do its work
 2. Send a completion message after the text is printed

PE: Failure Scenario

The server crashes, subsequently recovers and announces that to all clients

PE: Server Events Ordering

- Server events can occur in six different orderings:

Ordering	Description
$M \rightarrow P \rightarrow C$	A crash occurs after sending the completion message and printing the text
$M \rightarrow C(\rightarrow P)$	A crash occurs after sending the completion message, but before the text is printed
$P \rightarrow M \rightarrow C$	A crash occurs after printing the text and sending the completion message
$P \rightarrow C(\rightarrow M)$	The text is printed, after which a crash occurs before the completion message is sent
$C(\rightarrow P \rightarrow M)$	A crash happens before the server could do anything
$C(\rightarrow M \rightarrow P)$	A crash happens before the server could do anything

PE: Client Reissue Strategies

- After the crash of the server, the client does not know whether its request to print some text was carried out or not
- The client has a choice between 4 strategies:

Reissue Strategy	Description
Never	Never reissue a request, at the risk that the text will not be printed
Always	Always reissue a request, potentially leading to the text being printed twice
Reissue When Not ACKed	Reissue a request only if it did not yet receive an ACK that its print request had been delivered to the server
Reissue When ACKed	Reissue a request only if it has received an ACK for the print request

PE: Summary and Conclusion

- In summary, the different combinations of client and server strategies are as follows (*OK*= Text is printed once, *DUP*= Text is printed twice, and *ZERO*= Text is not printed):

Reissue Strategy	Strategy $M \rightarrow P$			Strategy $P \rightarrow M$		
	MPC	MC(P)	C(MP)	PMC	PC(M)	C(PM)
Always	DUP	OK	OK	DUP	DUP	OK
Never	OK	ZERO	ZERO	OK	OK	ZERO
Only when ACKed	DUP	OK	ZERO	DUP	OK	ZERO
Only when NOT ACKed	OK	ZERO	OK	OK	DUP	OK

PE: Summary and Conclusion

- In summary, the different combinations of client and server strategies are as follows (*OK= Text is printed once, DUP= Text is printed twice, and ZERO= Text is not printed*):

- There is NO combination of a client strategy and a server strategy that will work correctly under all possible event sequences
- The client can never know whether the server crashed just before or after having the text printed

Classes of Failures in Request-Reply Communication

- There are 5 different classes of failures that can occur in request-reply systems:
 1. The client is unable to locate the server
 2. The request message from the client to the server is lost
 3. The server crashes after receiving a request
 4. The reply message from the server to the client is lost
 5. The client crashes after sending a request

Possible Solution (1)

- The *doOperation* can start a *timer* when sending the request message
- If the timer expires before a reply or an ACK comes back, the message is sent again
- Considerations:
 - For that to happen, the client's request should be *idempotent*

Possible Solution (2)

- What if the client's request is not idempotent?
 - Have the client assign each request a sequence number
 - Have the server keep track of the most recently received sequence number from each client
 - The server can then tell the difference between an original request and a retransmission one and can refuse to carry out any request a second time

Classes of Failures in Request-Reply Communication

- There are 5 different classes of failures that can occur in request-reply systems:
 1. The client is unable to locate the server
 2. The request message from the client to the server is lost
 3. The server crashes after receiving a request
 4. The reply message from the server to the client is lost
 5. The client crashes after sending a request

Orphans

- A client might crash while the server is performing a corresponding computation
 - Such an unwanted computation is called an [orphan](#) (as there is no parent waiting for it after done)
- Orphans can cause a variety of problems that can interfere with the normal operation of the system:
 - They waste CPU cycles
 - They might lock up files and tie up valuable resources
 - If the client reboots, does the request again, and then an orphan reply comes back immediately afterwards, a confusion might occur


Possible Solutions [Nelson 1981]

- **S1: Extermination:** Use logging to explicitly kill off an orphan after a client reboot
- **S2: Reincarnation:** Use broadcasting to kill all remote computations on a client's behalf after rebooting and getting a new *epoch number*
- **S3: Gentle Reincarnation:** After an epoch broadcast comes in, a machine kills a computation only if its owner cannot be located anywhere
- **S4: Expiration:** each remote invocation is given a standard amount of time to fulfill the job

Orphan elimination is discussed in more detail by **Panzieri and Shrivastave (1988)**


The Two Generals Problem

- Two generals
 - lead divisions of an army
 - each camped on the mountains
 - on the two sides of an enemy-occupied valley
 - only communicate via messengers
- We need a scheme for the generals to agree on a common attack time
 - an attack by only one division would be disastrous
- Ideas?

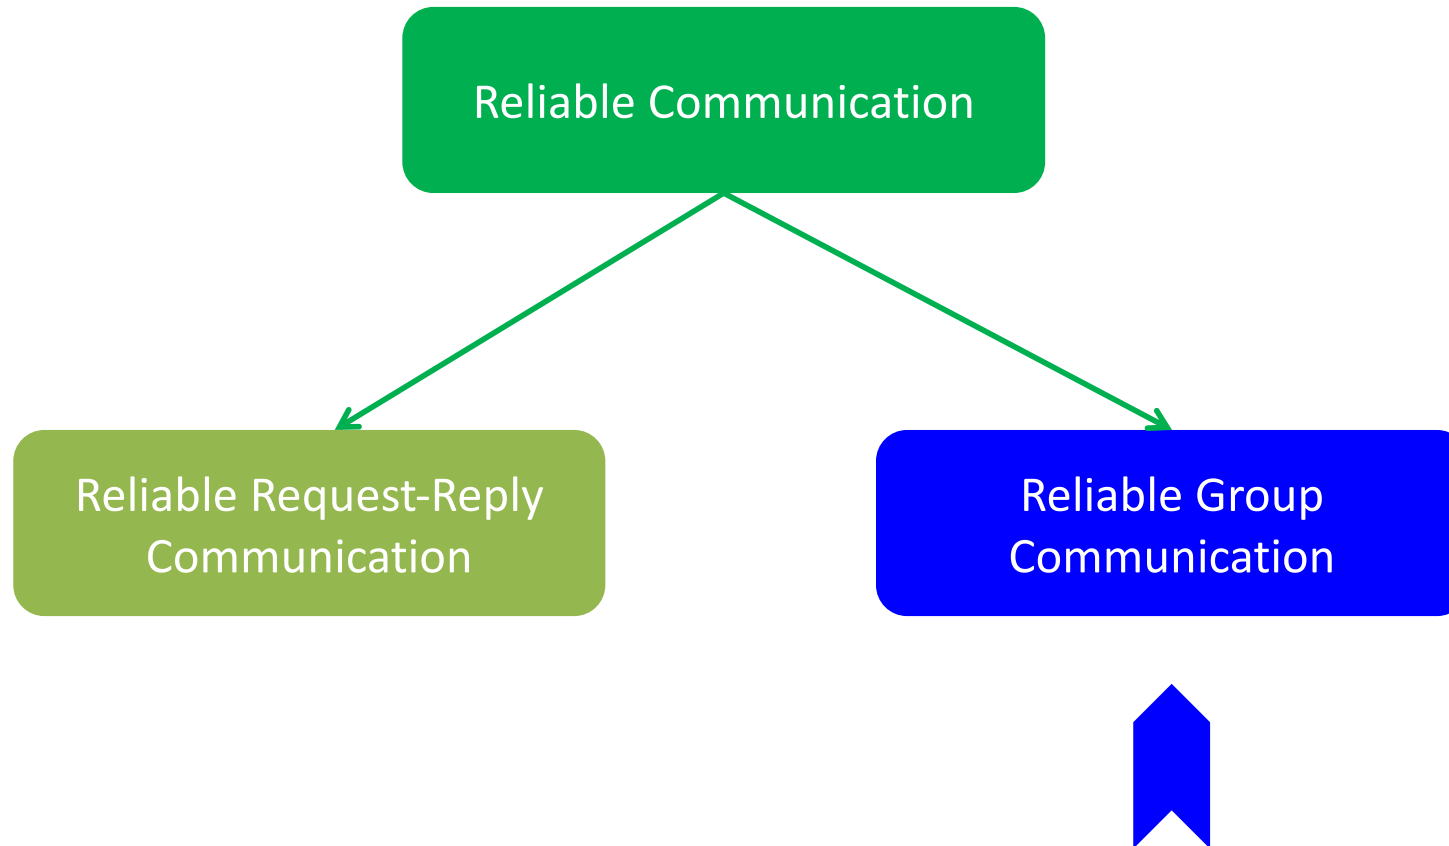


The Two Generals Problem

- Messengers are totally reliable, but may need an arbitrary amount of time to cross the valley (they may even be captured and never arrive)
 - G1 decides on T,
 - sends a messenger to tell G2
 - G2 acknowledges receipt of the attack time T
 - G2, unsure whether G1 got the ack (without which he would not attack), will need an ack of the ack!
- This can go on forever, without either being sure




Reliable Communication



Reliable Group Communication

- As we considered reliable request-reply communication, we need also to consider reliable multicasting services



- E.g., Election algorithms use multicasting schemes

Reliable Group Communication

- A Basic Reliable-Multicasting Scheme
- Scalability in Reliable Multicasting
- Atomic Multicast

Reliable Group Communication

- A Basic Reliable-Multicasting Scheme
- Scalability in Reliable Multicasting
- Atomic Multicast

Reliable Multicasting

- Reliable multicasting indicates that a message that is sent to a process group should be delivered to each member of that group
- A distinction should be made between:
 - Reliable communication in the presence of faulty processes
 - Reliable communication when processes are assumed to operate correctly
- In the presence of faulty processes, multicasting is considered to be reliable when it can be guaranteed that all non-faulty group members receive the message

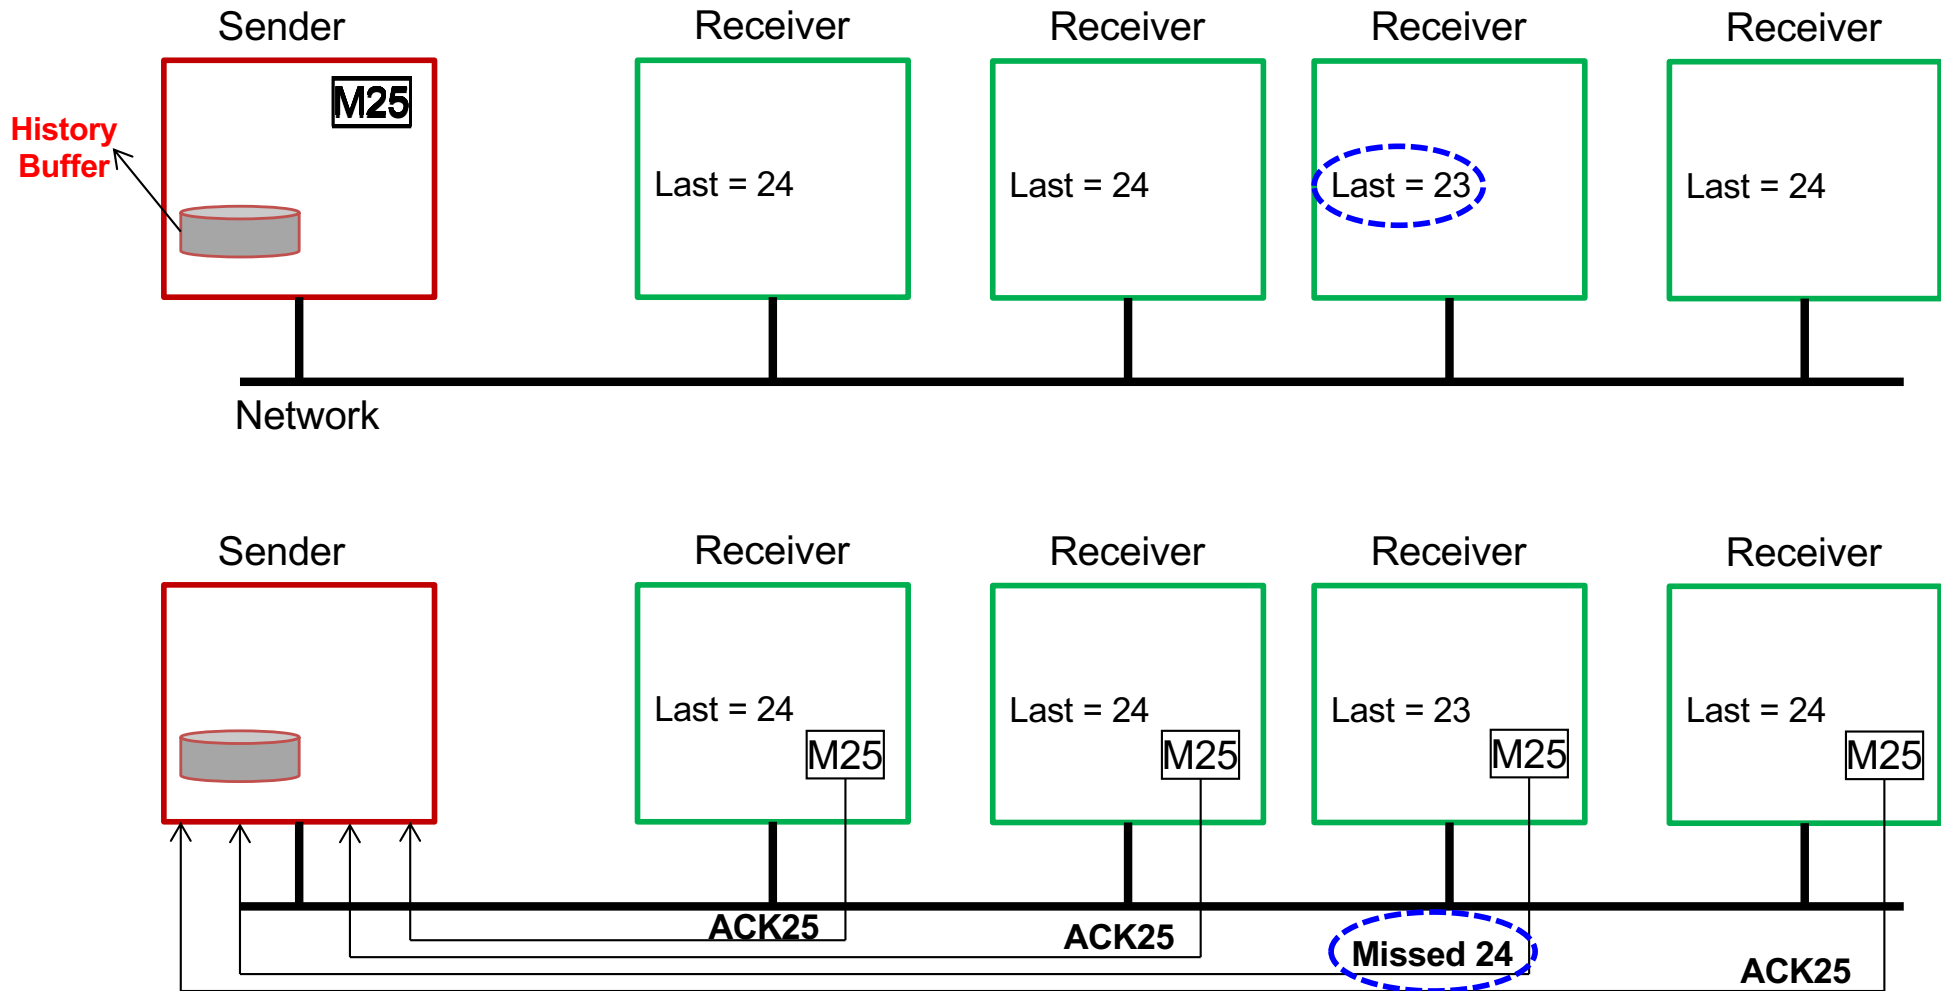
Basic Reliable Multicasting Questions

- What happens if during communication (i.e., a message is being delivered) a process P joins a group?
 - Should P also receive the message?
- What happens if a (sending) process crashes during communication?
- What about message ordering?

Reliable Multicasting with Feedback Messages

- Example: Node S multicasts a message to multiple receivers
- Problem: message may be lost part way
 - delivered to some, but not to all, of the intended receivers
- Assumption: messages are received in the same order as they are sent

Reliable Multicasting with Feedback Messages



An extensive and detailed survey of total-order broadcasts can be found in Defago et al. (2004)

Reliable Group Communication

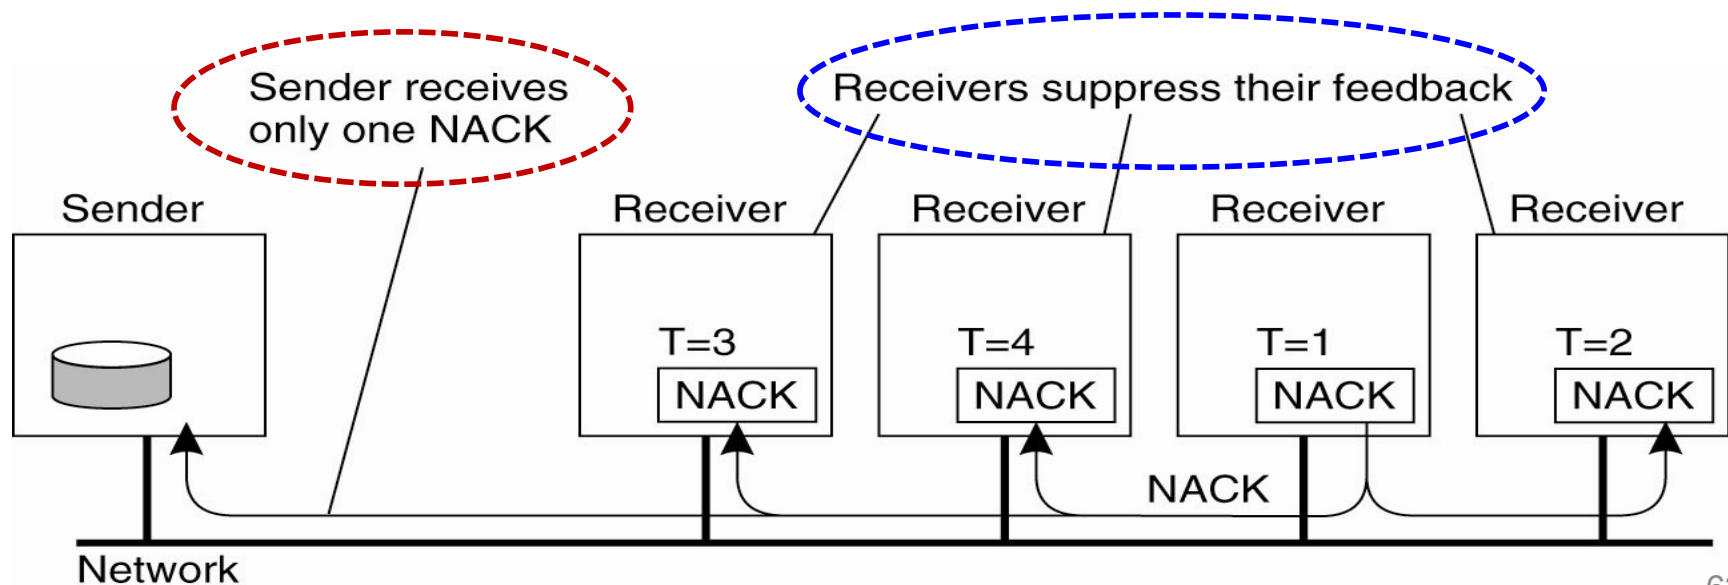
- A Basic Reliable-Multicasting Scheme
- Scalability in Reliable Multicasting
- Atomic Multicast

Scalability Issues with a Feedback-Based Scheme

- If there are N receivers in a multicasting process, the sender must be prepared to accept at least N ACKs
- This might cause a feedback implosion
 - Instead, we can let a receiver return only a NACK
- Limitations:
 - No hard guarantees can be given that a feedback implosion will not happen
 - It is not clear for how long the sender should keep a message in its history buffer


Nonhierarchical Feedback Control

- Control the number of NACKs sent back to the sender?
 - A NACK is sent to all the group members after some random delay
 - A group member suppresses its own feedback concerning a missing message after receiving a NACK feedback about the same message




Hierarchical Feedback Control

- Add structure (hierarchies)
 - Reduce feedback further
- Partition of receivers to subgroups
 - Organize as tree



Hierarchical Feedback Control

- The subgroup containing the sender *S* forms the root of the tree
- Within a subgroup, any reliable multicasting scheme can be used
- Each subgroup appoints a local coordinator *C* responsible for handling retransmission requests in its subgroup
- If *C* misses a message *m*, it asks the *C* of the parent subgroup to retransmit *m*



Next time:

- Continue with Reliable Group Communication
 - Atomic Multicast
- And much more
 - Atomicity and distributed commit protocols
 - Recovery from failures

Questions?

In part, inspired from / based on slides from

- Mohammad Hammoud
- Muyuan Wang
- Philippas Tsigas